



Crash工具基本使用及实战分享

目录

CONTENT

01

Crash工具基本使用

02

一个简单例子

03

现网问题实战

Crash工具基本使用

Crash是一个用于分析内核转储文件的分析工具

https://crash-utility.github.io/crash_whitepaper.html

启动命令：

```
crash vmlinux vmcore
```

vmlinux：未压缩的内核映像文件vmlinux，包含调试信息

vmcore：kdump生成的内核转储文件，如果此项不指定，默认分析实时系统内存

Crash工具基本使用

Crash基本输出

DTAE表示panic的时间

UPTIME是panic前已运行时间

TASKS表示发生panic时系统中进程总数

PANIC表示系统发生panic的原因及相应信息

PID是发生panic的任务号

COMMAND表示发生panic的任务正在执行的操作

TASK与THREAD_INFO表示发生panic的任务描述符及thread_info结构体的地址

CPU是发生panic的核

STATE表示发生panic的任务当时的状态

```
[root@localhost 127.0.0.1-2021-04-14-22:17:25]# crash vmlinux vmcore
KERNEL: /home/linux-rh-3-10/vmlinux
DUMPFILE: vmcore [PARTIAL DUMP]
CPUS: 16
DATE: Wed Apr 14 22:17:20 2021
UPTIME: 12:00:53
LOAD AVERAGE: 0.11, 0.06, 0.04
TASKS: 283
NODENAME: localhost.localdomain
RELEASE: 3.10.0-327.77.60.61.x86_64+
VERSION: #65 SMP Wed Jan 20 20:46:06 EST 2021
MACHINE: x86_64 (2593 Mhz)
MEMORY: 8 GB
PANIC: "BUG: unable to handle kernel NULL pointer dereference at 0000000000000804"
PID: 3754
COMMAND: "systemd-udev"
TASK: ffff880034ee8b90 [THREAD_INFO: ffff8800b90c0000]
CPU: 3
STATE: TASK_RUNNING (PANIC)

crash>
```

Crash工具基本使用

Crash常用命令

help 获取帮助信息，help 具体命令，如help bt, 可以查看详细使用帮助

```
crash> help
```

*	files	mach	repeat	timer
alias	foreach	mod	runq	tree
ascii	fuser	mount	search	union
bt	gdb	net	set	vm
btcp	help	p	sig	vtop
dev	ipcs	ps	struct	waitq
dis	irq	pte	swap	whatis
eval	kmem	ptob	sym	wr
exit	list	ptov	sys	q
extend	log	rd	task	

```
crash> help bt
```

```
NAME
  bt - backtrace
```

SYNOPSIS

```
bt [-a|-c cpu(s)|-g|-r|-t|-T|-l|-e|-E|-f|-F|-o|-O] [-R ref] [-s [-x|d]]
    [-I ip] [-S sp] [pid | task]
```

DESCRIPTION

Display a kernel stack backtrace. If no arguments are given, the stack trace of the current context will be displayed.

- a displays the stack traces of the active task on each CPU. (only applicable to crash dumps)
- A same as -a, but also displays vector registers (S390X only).
- c cpu display the stack trace of the active task on one or more CPUs, which can be specified using the format "3", "1,8,9", "1-23", or "1,8,9-14". (only applicable to crash dumps)
- g displays the stack traces of all threads in the thread group of the target task; the thread group leader will be displayed first.
- r display raw stack data, consisting of a memory dump of the two pages of memory containing the task_union structure.
- t display all text symbols found from the last known stack location to the top of the stack. (helpful if the back trace fails)
- T display all text symbols found from just above the task_struct or thread_info to the top of the stack. (helpful if the back trace fails or the -t option starts too high in the process stack).
- l show file and line number of each stack trace text location.

Crash工具基本使用

Crash常用命令

bt 显示函数调用栈， 可以显示所有CPU或指定CPU的栈， 或者指定pid

bt 显示当前cpu栈

bt -a 显示所有cpu栈

bt -f 显示所有堆栈

bt -l 显示堆栈tarce的文件与行号

```
crash> bt
PID: 3754 TASK: ffff880034ee8b90 CPU: 3 COMMAND: "systemd-udevd"
#0 [ffff8800b90c3840] machine_kexec at ffffffff81055deb
#1 [ffff8800b90c38a0] crash_kexec at ffffffff810fd992
#2 [ffff8800b90c3970] oops_end at ffffffff81660748
#3 [ffff8800b90c3998] no_context at ffffffff8164feb4
#4 [ffff8800b90c39e8] _bad_area_nosemaphore at ffffffff8164fff4
#5 [ffff8800b90c3a30] bad_area at ffffffff8165026e
#6 [ffff8800b90c3a58] _do_page_fault at ffffffff816636c9
#7 [ffff8800b90c3ab8] trace_do_page_fault at ffffffff81663816
#8 [ffff8800b90c3af8] do_async_page_fault at ffffffff81662ee9
#9 [ffff8800b90c3b10] async_page_fault at ffffffff8165f6e8
[exception RIP: tun_get_drvinfo+55]
RIP: ffffffff8055c2c7 RSP: ffff8800b90c3bc0 RFLAGS: 00010206
RAX: 0000000000000003 RBX: ffff8800b90c3bdc RCX: 0000000000000000
RDX: 0000000000000003 RSI: ffffffff8055f0e6 RDI: ffff8800b90c3c03
RBP: ffff8800b90c3bc8 RB: 0000000000000000 R9: ffff3e465a9e7574
R10: 00007fd033b4cc30 R11: 0000000000000246 R12: ffffffff8055f940
R13: ffff8800b90c3bdc R14: 00007fff3deba7d0 R15: ffffffff81a58fc0
ORIG_RAX: ffffffff8165f6e8 CS: 0010 SS: 0018
#10 [ffff8800b90c3bd0] ethtool_get_drvinfo at ffffffff8154da64
#11 [ffff8800b90c3cd0] dev_ethtool at ffffffff8154f486
#12 [ffff8800b90c3da8] dev_ioctl at ffffffff8155f9ef
#13 [ffff8800b90c3e38] sock_do_ioctl at ffffffff8152c18d
#14 [ffff8800b90c3e60] sock_ioctl at ffffffff8152c8a8
#15 [ffff8800b90c3e90] do_vfs_ioctl at ffffffff81202988
#16 [ffff8800b90c3f00] sys_ioctl at ffffffff81202c11
#17 [ffff8800b90c3f50] system_call_fastpath at ffffffff81668155
RIP: 00007fd033b5507 RSP: 00007fff3deba788 RFLAGS: 00010246
RAX: 0000000000000010 RBX: 000055e501678e98 RCX: ffffffff816680a1
RDX: 00007fff3deba7a0 RSI: 0000000000008946 RDI: 0000000000000008
RBP: 00007fff3deba7a0 RB: 00000000000000ff R9: 0000000000000007
R10: 00007fd033b4cc30 R11: 0000000000000246 R12: 00007fff3deba7d0
R13: 00007fff3deba970 R14: 000055e5016800 R15: 0000000000000000
ORIG_RAX: 0000000000000010 CS: 0033 SS: 002b
```

```
crash> bt -a
PID: 0 TASK: ffffffff81982440 CPU: 0 COMMAND: "swapper/0"
#0 [ffff880237208e58] crash_nmi_callback at ffffffff81048768
#1 [ffff880237208e68] nmi_handle at ffffffff816608df
#2 [ffff880237208eb0] do_nmi at ffffffff81660a54
#3 [ffff880237208ef0] end_repeat_nmi at ffffffff8165fcec
[exception RIP: native_safe_halt+11]
RIP: ffffffff8105cf2b RSP: ffffffff8196fea0 RFLAGS: 00000246
RAX: ffffffff8101ed50 RBX: ffffffff8196c000 RCX: 0100000000000000
RDX: 0000000000000000 RSI: 0000000000000000 RDI: 0000000000000046
RBP: ffffffff8196fea0 RB: 0000000000000000 R9: 001b615238000000
R10: 0000000000000000 R11: ffff880233a11c00 R12: 0000000000000000
R13: ffffffff8196c000 R14: ffffffff8196c000 R15: ffffffff8196ffa8
ORIG_RAX: ffffffff8165fcec CS: 0010 SS: 0018
--- <NMI exception stack> ---
#4 [fffffff8196fea0] native_safe_halt at ffffffff8105cf2b
#5 [fffffff8196fea8] default_idle at ffffffff8101ed6f
#6 [fffffff8196fec8] arch_cpu_idle at ffffffff8101fab0
#7 [fffffff8196fed8] cpu_startup_entry at ffffffff810df7d5
#8 [fffffff8196ff30] rest_init at ffffffff816467ac
#9 [fffffff8196ff40] start_kernel at ffffffff81ac50ed
#10 [fffffff8196ff88] x86_64_start_reservations at ffffffff81ac466c
#11 [fffffff8196ff98] x86_64_start_kernel at ffffffff81ac47c0
PID: 0 TASK: ffff880233f26810 CPU: 1 COMMAND: "swapper/1"
#0 [ffff880237248e58] crash_nmi_callback at ffffffff81048768
#1 [ffff880237248e68] nmi_handle at ffffffff816608df
#2 [ffff880237248eb0] do_nmi at ffffffff81660a54
#3 [ffff880237248ef0] end_repeat_nmi at ffffffff8165fcec
[exception RIP: native_safe_halt+11]
RIP: ffffffff8105cf2b RSP: ffff880233f4be98 RFLAGS: 00000246
RAX: ffffffff8101ed50 RBX: ffff880233f48000 RCX: 0100000000000000
RDX: 0000000000000000 RSI: 0000000000000000 RDI: 0000000000000046
RBP: ffff880233f4be98 RB: 0000000000000000 R9: 001b615238000000
R10: 0000000000000000 R11: 0000000000000005 R12: 0000000000000001
R13: ffff880233f48000 R14: ffff880233f48000 R15: 0000000000000000
ORIG_RAX: ffffffff8165fcec CS: 0010 SS: 0018
```

```
crash> bt -f
PID: 3754 TASK: ffff880034ee8b90 CPU: 3 COMMAND: "systemd-udevd"
#0 [ffff8800b90c3840] machine_kexec at ffffffff81055deb
ffff8800b90c3848: 00008800b90c38a8 ffff880000000000
ffff8800b90c3958: 0000000013001000 ffff880013001000
ffff8800b90c3968: 0000000013000000 ffff8800b90c3888
ffff8800b90c3878: 00000000b7df7224 ffff8800b90c3b18
ffff8800b90c3888: ffff8800b90c38a8 0000000000000246
ffff8800b90c3898: ffff8800b90c3968 ffffffff810fd992
#1 [ffff8800b90c38a0] crash_kexec at ffffffff810fd992
ffff8800b90c38a8: ffffffff81a58fcd 0000/fff3deba/d0
ffff8800b90c38b8: ffff8800b90c3bdc ffffffff8055f940
ffff8800b90c3928: ffff8800b90c3bc8 ffff8800b90c3bdc
ffff8800b90c38d8: 0000000000000246 00007fd033b4cc30
ffff8800b90c38e8: ffff3e465a9e7574 0000000000000000
ffff8800b90c38f8: 0000000000000003 0000000000000000
ffff8800b90c3908: 0000000000000003 ffffffff8055f0e6
ffff8800b90c3918: ffff8800b90c3c03 ffffffff816636c9
ffff8800b90c3928: ffffffff8055c2c7 0000000000000010
ffff8800b90c3938: 0000000000010206 ffff8800b90c3bdc
ffff8800b90c3948: 0000000000000018 00000000b7df7224
ffff8800b90c3958: 0000000000000009 ffff8800b90c3b18
ffff8800b90c3968: ffff8800b90c3990 ffffffff81660748
#2 [ffff8800b90c3970] oops_end at ffffffff81660748
ffff8800b90c3978: ffff8800b90c3b18 0000000000000009
ffff8800b90c3988: 0000000000000246 ffff8800b90c39e0
ffff8800b90c3998: ffffffff8164feb4
#3 [ffff8800b90c3998] no_context at ffffffff8164feb4
ffff8800b90c39a0: 0003000100000000 ffff88023fd0e008
ffff8800b90c39b0: 00000000b7df7224 0000000000000000
ffff8800b90c39c0: ffff8800b90c3b18 0000000000000004
ffff8800b90c39d0: ffff880034ee8b90 0000000000030001
ffff8800b90c39e0: ffff8800b90c3a28 ffffffff8164fff4
#4 [ffff8800b90c39e8] _bad_area_nosemaphore at ffffffff8164fff4
ffff8800b90c39f0: 000000000000001c 0000000000000000
ffff8800b90c3a00: ffff8800b90c3b18 0000000000000008
ffff8800b90c3a10: 0000000000000004 000000000000002b
ffff8800b90c3a20: ffff88022bf52580 ffff8800b90c3a50
ffff8800b90c3a30: ffffffff8165026e
```

```
crash> bt -l
PID: 3754 TASK: ffff880034ee8b90 CPU: 3 COMMAND: "systemd-udevd"
#0 [ffff8800b90c3840] machine_kexec at ffffffff81055deb
/home/linux-rh-3-10/arch/x86/kernel/machine_kexec_64.c: 327
#1 [ffff8800b90c38a0] crash_kexec at ffffffff810fd992
/home/linux-rh-3-10/kernel/kexec.c: 1524
#2 [ffff8800b90c3970] oops_end at ffffffff81660748
/home/linux-rh-3-10/arch/x86/kernel/dumpstack.c: 224
#3 [ffff8800b90c3998] no_context at ffffffff8164feb4
/home/linux-rh-3-10/arch/x86/mm/fault.c: 714
#4 [ffff8800b90c39e8] _bad_area_nosemaphore at ffffffff8164fff4a
/home/linux-rh-3-10/arch/x86/mm/fault.c: 793
#5 [ffff8800b90c3a30] bad_area at ffffffff8165026e
/home/linux-rh-3-10/arch/x86/mm/fault.c: 822
#6 [ffff8800b90c3a58] _do_page_fault at ffffffff816636c9
/home/linux-rh-3-10/arch/x86/mm/fault.c: 1163
#7 [ffff8800b90c3ab8] trace_do_page_fault at ffffffff81663816
/home/linux-rh-3-10/arch/x86/mm/fault.c: 1263
#8 [ffff8800b90c3af8] do_async_page_fault at ffffffff81662ee9
/home/linux-rh-3-10/arch/x86/kernel/kvm.c: 288
#9 [ffff8800b90c3b10] async_page_fault at ffffffff8165f6e8
/home/linux-rh-3-10/arch/x86/kernel/entry_64.S: 1390
[exception RIP: tun_get_drvinfo+55]
RIP: ffffffff8055c2c7 RSP: ffff8800b90c3bc0 RFLAGS: 00010206
RAX: 0000000000000003 RBX: ffff8800b90c3bdc RCX: 0000000000000000
RDX: 0000000000000003 RSI: ffffffff8055f0e6 RDI: ffff8800b90c3c03
RBP: ffff8800b90c3bc8 RB: 0000000000000000 R9: ffff3e465a9e7574
R10: 00007fd033b4cc30 R11: 0000000000000246 R12: ffffffff8055f940
R13: ffff8800b90c3bdc R14: 00007fff3deba7d0 R15: ffffffff81a58fcd
ORIG_RAX: ffffffff8165f6e8 CS: 0010 SS: 0018
/home/linux-rh-3-10/drivers/net/tun.c: 2401
#10 [ffff8800b90c3bd0] ethtool_get_drvinfo at ffffffff8154da64
/home/linux-rh-3-10/net/core/ethtool.c: 396
#11 [ffff8800b90c3cd0] dev_ethtool at ffffffff8154f486
/home/linux-rh-3-10/net/core/ethtool.c: 1687
```



Crash工具基本使用

Crash常用命令

#0 代表callTrace层级

exception RIP 异常指令

RAX RBX RCX 寄存器值

#10 [ffff8800b90c3bd0] ethtool_get_drvinfo at ffffffff8154da64

系统在使用线性地址为ffff8800b90c3bd0处的空间时，在ethtool_get_drvinfo

函数中运行，从ffffffff8154da64处开始调用下一层函数tun_get_drvinfo

```
#9 [ffff8800b90c3b10] async_page_fault at ffffffff8165f6e8
[exception RIP: tun_get_drvinfo+55]
RIP: ffffffff8055c2c7 RSP: ffff8800b90c3bc0 RFLAGS: 00010206
RAX: 0000000000000003 RBX: ffff8800b90c3bdc RCX: 0000000000000000
RDX: 0000000000000003 RSI: ffffffff8055f0e6 RDI: ffff8800b90c3c03
RBP: ffff8800b90c3bc8 R8: 0000000000000000 R9: ffff3e465a9e7574
R10: 00007fd033b4cc30 R11: 0000000000000246 R12: ffffffff8055f940
R13: ffff8800b90c3bdc R14: 00007fff3deba7d0 R15: ffffffff81a58fc0
ORIG_RAX: ffffffff81a58fc0 CS: 0010 SS: 0018
ffff8800b90c3b18: ffffffff81a58fc0 00007fff3deba7d0
ffff8800b90c3b28: ffff8800b90c3bdc ffffffff8055f940
ffff8800b90c3b38: ffff8800b90c3bc8 ffff8800b90c3bdc
ffff8800b90c3b48: 0000000000000246 00007fd033b4cc30
ffff8800b90c3b58: ffff3e465a9e7574 0000000000000000
ffff8800b90c3b68: 0000000000000003 0000000000000000
ffff8800b90c3b78: 0000000000000003 ffffffff8055f0e6
ffff8800b90c3b88: ffff8800b90c3c03 ffffffff81a58fc0
ffff8800b90c3b98: ffffffff8055c2c7 0000000000000010
ffff8800b90c3ba8: 0000000000010206 ffff8800b90c3bc0
ffff8800b90c3bb8: 0000000000000018 ffff8800357d4000
ffff8800b90c3bc8: ffff8800b90c3cc8 ffffffff8154da64
#10 [ffff8800b90c3bd0] ethtool_get_drvinfo at ffffffff8154da64
ffff8800b90c3bd8: 0000000335c9b500 00000000006e7574
ffff8800b90c3be8: 0000000000000000 0000000000000000
ffff8800b90c3bf8: 0000000000000000 0000000000362e31
ffff8800b90c3c08: 0000000000000000 0000000000000000
ffff8800b90c3c18: 0000000000000000 0000000000000000
ffff8800b90c3c28: 0000000000000000 0000000000000000
```

```
crash> bt
PID: 3754 TASK: ffff880034ee8b90 CPU: 3 COMMAND: "systemd-udev"
#0 [ffff8800b90c3840] machine_kexec at ffffffff81055deb
#1 [ffff8800b90c38a0] crash_kexec at ffffffff810fd992
#2 [ffff8800b90c3970] oops_end at ffffffff81660748
#3 [ffff8800b90c3998] no_context at ffffffff8164feb4
#4 [ffff8800b90c39e8] __bad_area_nosemaphore at ffffffff8164ff4a
#5 [ffff8800b90c3a30] bad_area at ffffffff8165026e
#6 [ffff8800b90c3a58] __do_page_fault at ffffffff816636c9
#7 [ffff8800b90c3ab8] trace_do_page_fault at ffffffff81663816
#8 [ffff8800b90c3af8] do_async_page_fault at ffffffff81662ee9
#9 [ffff8800b90c3b10] async_page_fault at ffffffff8165f6e8
[exception RIP: tun_get_drvinfo+55]
RIP: ffffffff8055c2c7 RSP: ffff8800b90c3bc0 RFLAGS: 00010206
RAX: 0000000000000003 RBX: ffff8800b90c3bdc RCX: 0000000000000000
RDX: 0000000000000003 RSI: ffffffff8055f0e6 RDI: ffff8800b90c3c03
RBP: ffff8800b90c3bc8 R8: 0000000000000000 R9: ffff3e465a9e7574
R10: 00007fd033b4cc30 R11: 0000000000000246 R12: ffffffff8055f940
R13: ffff8800b90c3bdc R14: 00007fff3deba7d0 R15: ffffffff81a58fc0
ORIG_RAX: ffffffff81a58fc0 CS: 0010 SS: 0018
#10 [ffff8800b90c3bd0] ethtool_get_drvinfo at ffffffff8154da64
#11 [ffff8800b90c3cd0] dev_ethtool at ffffffff8154f486
#12 [ffff8800b90c3da8] dev_ioctl at ffffffff8155f9ef
#13 [ffff8800b90c3e38] sock_do_ioctl at ffffffff8152c18d
#14 [ffff8800b90c3e60] sock_ioctl at ffffffff8152c8a8
#15 [ffff8800b90c3e90] do_vfs_ioctl at ffffffff81202988
#16 [ffff8800b90c3f00] sys_ioctl at ffffffff81202c11
#17 [ffff8800b90c3f50] system_call_fastpath at ffffffff81668155
RIP: 00007fd033bb5507 RSP: 00007fff3deba788 RFLAGS: 00010246
RAX: 0000000000000010 RBX: 0000556e501678e8 RCX: ffffffff816680a1
RDX: 00007fff3deba7a0 RSI: 0000000000000894 RDI: 0000000000000008
RBP: 00007fff3deba7a0 R8: 0000000000000fff R9: 0000000000000007
R10: 00007fd033b4cc30 R11: 0000000000000246 R12: 00007fff3deba7d0
R13: 00007fff3deba970 R14: 0000556e50168600 R15: 0000000000000000
ORIG_RAX: 0000000000000010 CS: 0033 SS: 002b
```


Crash工具基本使用

Crash常用命令

log 显示系统消息缓存区

```
crash> log
[ 0.000000] Initializing cgroup subsys cpuset
[ 0.000000] Initializing cgroup subsys cpu
[ 0.000000] Initializing cgroup subsys cpuacct
[ 0.000000] Linux version 3.10.0-327.77.60.61.x86_64+ (root@localhost.localdomain) (gcc version 4.8.3)
[ 0.000000] Command line: BOOT_IMAGE=/vmlinuz-3.10.0-327.77.60.61.x86_64+ root=/dev/mapper/rhel-t
on
[ 0.000000] e820: BIOS-provided physical RAM map:
[ 0.000000] BIOS-e820: [mem 0x0000000000000000-0x000000000009fbff] usable
[ 0.000000] BIOS-e820: [mem 0x000000000009fc00-0x00000000000fffff] reserved
[ 0.000000] BIOS-e820: [mem 0x00000000000ffff0-0x00000000000fffff] reserved
[ 0.000000] BIOS-e820: [mem 0x0000000001000000-0x00000000000bffff] usable
[ 0.000000] BIOS-e820: [mem 0x00000000000bffc000-0x00000000000bffff] reserved
[ 0.000000] BIOS-e820: [mem 0x00000000000fffc000-0x00000000000fefffff] reserved
[ 0.000000] BIOS-e820: [mem 0x00000000000fffc000-0x00000000000fffc000] reserved
[ 0.000000] BIOS-e820: [mem 0x000000000100000000-0x00000000023fffff] usable
[ 0.000000] NX (Execute Disable) protection: active
[ 0.000000] SMBIOS 2.4 present.
[ 0.000000] DMI: Red Hat KVM, BIOS 0.5.1 01/01/2011
[ 0.000000] Hypervisor detected: KVM
[ 0.000000] e820: update [mem 0x00000000-0x00000fff] usable ==> reserved
[ 0.000000] e820: remove [mem 0x00000000-0x000fffff] usable
```

```
[43254.289274] BUG: unable to handle kernel NULL pointer dereference at 0000000000000804
[43254.292156] IP: [<ffffffffffa055c2c7>] tun_get_drvinfo+0x37/0x80 [tun]
[43254.293008] PGD 0
[43254.293297] Oops: 0000 [#1] SMP
[43254.293762] Modules linked in: tun(OE) ip6t_rpfilter ipt_REJECT ip6t_REJECT xt_conntrack ip_set nfnetlink ebtabel_nat ebtabel
ip6table_mangle ip6table_security ip6table_raw iptable_nat nf_conntrack_ipv4 nf_defrag_ipv4 nf_nat_ipv4 nf_nat nf_conntrack ip
p6_tables iptable_filter ext4 jbd2 mbcache ppdev i2c_piix4 virtio_balloon parport_pc pcspkr parport sg nfsd auth_rpcgss nfs_ac
eric cdrom crc10dif_common ata_generic pata_acpi cirrus syscopyarea sysfillrect sysimgblt drm_kms_helper ttm drm 8139too virt
[43254.303884] virtio dm_mirror dm_region_hash dm_log dm_mod
[43254.304526] CPU: 3 PID: 3754 Comm: systemd-udevd Tainted: G OE ---- - - - - - 3.10.0-327.77.60.61.x86_64+ #65
[43254.305872] Hardware name: Red Hat KVM, BIOS 0.5.1 01/01/2011
[43254.306599] task: ffff880034ee8b90 ti: ffff8800b90c0000 task.ti: ffff8800b90c0000
[43254.307534] RIP: 0010:[<ffffffffffa055c2c7>] [<ffffffffffa055c2c7>] tun_get_drvinfo+0x37/0x80 [tun]
[43254.308652] RSP: 0018:ffff8800b90c3bc0 EFLAGS: 00010206
[43254.309311] RAX: 0000000000000003 RBX: ffff8800b90c3bdc RCX: 0000000000000000
[43254.310201] RDX: 0000000000000003 RSI: ffffffffa055f0e6 RDI: ffff8800b90c3c03
[43254.311091] RBP: ffff8800b90c3bc8 R08: 0000000000000000 R09: ffff3e465a9e7574
[43254.311982] R10: 00007fd033b4cc30 R11: 0000000000000246 R12: ffffffffa055f940
[43254.312873] R13: ffff8800b90c3bdc R14: 00007fff3deba7d0 R15: ffffffffa055f940
[43254.313760] FS: 00007fd034f548c0(0000) GS: ffff8802372c0000(0000) knlGS: 0000000000000000
[43254.314766] CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
[43254.315489] CR2: 0000000000000804 CR3: 00000000357aa000 CR4: 00000000000006e0
[43254.316390] DR0: 0000000000000000 DR1: 0000000000000000 DR2: 0000000000000000
[43254.317271] DR3: 0000000000000000 DR6: 00000000ffff00ff DR7: 0000000000000400
[43254.318163] Call Trace:
[43254.318513] [<ffffffffff8154da64>] ethtool_get_drvinfo+0x84/0x1d0
[43254.319261] [<ffffffffff8154f486>] dev_ethtool+0xa16/0x1b30
[43254.320015] [<ffffffffffa0224dbc>] ? xfs_iunlock+0x11c/0x130 [xfs]
[43254.320796] [<ffffffffff81175efb>] ? unlock_page+0x2b/0x30
[43254.321478] [<ffffffffff815486dc>] ? dev_get_by_name_rcu+0x5c/0x80
[43254.322240] [<ffffffffff8155f9ef>] dev_ioctl+0x1cf/0x590
[43254.322908] [<ffffffffff811a535c>] ? handle_mm_fault+0x65c/0x1010
[43254.323663] [<ffffffffff8152c18d>] sock_do_ioctl+0x4d/0x60
[43254.324333] [<ffffffffff8152c8a8>] sock_ioctl+0x1f8/0x2d0
[43254.325011] [<ffffffffff81202988>] do_vfs_ioctl+0x2e8/0x4d0
[43254.325722] [<ffffffffff81663471>] ? __do_page_fault+0x171/0x430
[43254.326466] [<ffffffffff816680a1>] ? system_call_after_swapgs+0xae/0x146
[43254.327285] [<ffffffffff81202c11>] SyS_ioctl+0xa1/0xc0
[43254.327924] [<ffffffffff81668095>] ? system_call_after_swapgs+0xa2/0x146
[43254.328753] [<ffffffffff81668155>] system_call_fastpath+0x1c/0x21
[43254.329507] [<ffffffffff816680a1>] ? system_call_after_swapgs+0xae/0x146
[43254.330321] Code: 00 00 00 48 89 e5 53 48 89 f3 48 c7 c6 df 05 55 a0 e8 1e 4f db e0 48 8d 7b 24 ba 20 00 00 00 48 c7 c6 e3
15 48
[43254.333842] RIP [<ffffffffffa055c2c7>] tun_get_drvinfo+0x37/0x80 [tun]
[43254.334668] RSP <ffff8800b90c3bc0>
[43254.335109] CR2: 0000000000000804
```

	Value	
Bit	0	1
0	No page found	Invalid access
1	Read or Execute	Write
2	Kernel mode	User mode
3	Not instruction fetch	Instruction fetch



Crash工具基本使用

Crash常用命令

sym symbol与虚拟地址转换

```
crash> sym ethtool_get_drvinfo
ffffffff8154d9e0 (t) ethtool_get_drvinfo /home/linux-rh-3-10/net/core/ethtool.c: 375
crash>
crash> sym ffffffff8152c8a8
ffffffff8152c8a8 (t) sock_ioctl+504 /home/linux-rh-3-10/net/socket.c: 1097
crash> █
```

net 列出网络设备

```
crash> net
      NET_DEVICE      NAME      IP ADDRESS(ES)
ffff880233bf7000      lo        127.0.0.1
ffff88022b9f8000      ens3      9.84.160.187
ffff88022b9fb000      ens5
ffff8800357d4000      tun0
crash> █
```

net 显示ARP cache

```
crash> net -a
NEIGHBOUR      IP ADDRESS      HW TYPE      HW ADDRESS      DEVICE      STATE
ffff88022ac8ba00 9.84.101.101    ETHER        10:47:80:00:5e:1f ens3      STALE
ffff88022ac88200 9.84.0.62      ETHER        e4:c2:d1:f6:02:90 ens3      STALE
ffff88022ac89800 9.84.0.51      ETHER        20:3d:b2:46:ff:f0 ens3      STALE
ffff88022ac8bc00 9.84.0.14      ETHER        70:fd:45:5f:49:e2 ens3      STALE
ffff88022ac89200 127.0.0.1      UNKNOWN     00 00 00 00 00 00 lo        NOARP
ffff88022ac89e00 9.84.0.64      ETHER        70:fd:45:5f:42:92 ens3      STALE
ffff88022ac88c00 9.84.0.50      ETHER        68:8f:84:01:60:7e ens3      STALE
ffff88022ac8b000 9.84.0.23      ETHER        70:fd:45:eb:81:22 ens3      STALE
ffff88022befde00 9.84.0.1      ETHER        f4:1d:6b:87:4d:d2 ens3      REACHABLE
crash> █
```

net 列出所有sock

```
crash> net -s
PID: 3754  TASK: ffff880034ee8b90  CPU: 3  COMMAND: "systemd-udevd"
FD      SOCKET      SOCK      FAMILY:TYPE  SOURCE-PORT  DESTINATION-PORT
1 ffff8800351423c0 ffff88022ae96000 UNIX:STREAM
2 ffff8800351423c0 ffff88022ae96000 UNIX:STREAM
5 ffff8800351782c0 ffff88022af4a000 UNIX:DGRAM
8 ffff88023262a680 ffff88022a460000 INET:DGRAM    0.0.0.0-0 0.0.0.0-0
10 ffff880035178840 ffff88022af49800 UNIX:DGRAM
12 ffff8802325d3180 ffff8800b929c000 NETLINK/ROUTE:RAW
crash> █
```

rd 直接读内存

```
crash> rd ffff88022b9f8000 8
ffff88022b9f8000: 0000000033736e65 0000000000000000 ens3.....
ffff88022b9f8010: 0000000000000000 ffff880233bce100 .....3....
ffff88022b9f8020: 0000000000000000 0000000000000000 .....
ffff88022b9f8030: 0000000000000000 0000000000000000 .....
crash> █
```

Crash工具基本使用

Crash常用命令

set 获取crash的线程号

```
crash> set
PID: 3754
COMMAND: "systemd-udevd"
TASK: ffff880034ee8b90 [THREAD_INFO: ffff8800b90c0000]
CPU: 3
STATE: TASK_RUNNING (PANIC)
```

struct 解析结构体

```
crash> struct net_device ffff88022b9f8000
struct net_device {
    name = "ens3\000\000\000\000\000\000\000\000\000\000",
    name_hlist = {
        next = 0x0,
        pprev = 0xffff880233bce100
    },
    ifalias = 0x0,
    mem_end = 0.
```

```
crash> struct net_device ffff88022b9f8000 -o
struct net_device {
    [ffff88022b9f8000] char name[16];
    [ffff88022b9f8010] struct hlist_node name_hlist;
    [ffff88022b9f8020] char *ifalias;
    [ffff88022b9f8028] unsigned long mem_end;
    [ffff88022b9f8030] unsigned long mem_start;
    [ffff88022b9f8038] unsigned long base_addr;
```

dis 反汇编地址

```
crash> dis -l ethtool_get_drvinfo
/home/linux-rh-3-10/net/core/ethtool.c: 375
0xffffffff8154d9e0 <ethtool_get_drvinfo>:      nopl    0x0(%rax,%rax,1) [FTRACE NOP]
0xffffffff8154d9e5 <ethtool_get_drvinfo+5>:      push    %rbp
/home/linux-rh-3-10/net/core/ethtool.c: 379
0xffffffff8154d9e6 <ethtool_get_drvinfo+6>:      mov     $0xc4,%edx
/home/linux-rh-3-10/net/core/ethtool.c: 375
0xffffffff8154d9eb <ethtool_get_drvinfo+11>:     mov     %rsp,%rbp
0xffffffff8154d9ee <ethtool_get_drvinfo+14>:     push    %r14
0xffffffff8154d9f0 <ethtool_get_drvinfo+16>:     mov     %rsi,%r14
0xffffffff8154d9f3 <ethtool_get_drvinfo+19>:     push    %r13
/home/linux-rh-3-10/net/core/ethtool.c: 379
0xffffffff8154d9f5 <ethtool_get_drvinfo+21>:     lea     -0xec(%rbp),%r13
```

kmem -S 显示slab对象信息

```
crash> kmem -S
CACHE                                NAME                                OBJSIZE  ALLOCATED  TOTAL  SLABS  SSIZE
ffff8802322e6100 nf_conntrack_ffffffff81a58fc0 312      2          400      16      8k
CPU 0 KMEM CACHE CPU:
      ffffe8ffffa02d20
CPU 0 SLAB:
SLAB                                MEMORY                                NODE  TOTAL  ALLOCATED  FREE
ffffea0008b49980 ffff88022d266000 0      25      1      24
FREE / [ALLOCATED]
[ffff88022d266000]
ffff88022d266140 (cpu 0 cache)
ffff88022d266280 (cpu 0 cache)
```

Crash工具基本使用

Crash常用命令

p 查看全局变量

```
crash> p sysctl_tcp_rmem
sysctl_tcp_rmem = $7 =
{4096, 87380, 6291456}
```

dev命令可以显示系统中块设备和字符设备的信息

```
crash> dev
```

CHRDEV	NAME	CDEV	OPERATIONS
1	mem	ffff880233baba80	memory_fops
4	/dev/vc/0	fffffffff81f56000	console_fops
4	tty	ffff8802334a0000	tty_fops
4	ttyS	ffff88022bd40000	tty_fops
5	/dev/tty	fffffffff81f54f00	tty_fops
5	/dev/console	fffffffff81f54e80	console_fops

files命令可以显示发生panic的任务所打开的所有文件的信息

```
crash> files
PID: 3754 TASK: ffff880034ee8b90 CPU: 3 COMMAND: "systemd-udev"
ROOT: / CWD: /
```

FD	FILE	DENTRY	INODE	TYPE	PATH
0	ffff880233786700	ffff88023680a240	ffff880233a90868	CHR	/dev/null
1	ffff880233786800	ffff88022b056900	ffff8800351423f0	SOCK	UNIX
2	ffff880233786800	ffff88022b056900	ffff8800351423f0	SOCK	UNIX
3	ffff880232fed300	ffff8800b1e8bc80	ffff880236bd1770	UNKN	[signalfd]
4	ffff880232feda00	ffff8800b1e8bbc0	ffff880236bd1770	UNKN	[eventpoll]
5	ffff8800354eae00	ffff880236bffd40	ffff8800351782f0	SOCK	UNIX

Irq命令可以显示相应中断信息

```
crash> irq
```

IRQ	IRQ_DESC/ DATA	IRQACTION	NAME
0	ffff880236d14000	fffffffff819864c0	"timer"
1	ffff880236d14100	ffff880233ae8480	"i8042"
2	ffff880236d14200	(unused)	
3	ffff880236d14300	(unused)	
4	ffff880236d14400	ffff88022d5fbc80	"serial"
5	ffff880236d14500	(unused)	
6	ffff880236d14600	(unused)	
7	ffff880236d14700	(unused)	
8	ffff880236d14800	ffff880233a62f80	"rtc0"
9	ffff880236d14900	ffff880233b11f80	"acpi"
10	ffff880236d14a00	ffff88022975ec80	"ens3"
		ffff880034faff00	"ens5"

mount 用于显示挂载的文件系统信息

```
crash> mount
```

MOUNT	SUPERBLK	TYPE	DEVNAME	DIRNAME
ffff880233d88140	ffff880233d98800	rootfs	rootfs	/
ffff88022bf92000	ffff88022bfa0000	sysfs	sysfs	/sys
ffff88022bf92140	ffff880233d9b000	proc	proc	/proc
ffff88022bf92280	ffff880233a88000	devtmpfs	devtmpfs	/dev
ffff88022bf923c0	ffff88022be31000	securityfs	securityfs	/sys/kernel/security
ffff88022bf92500	ffff88022bfa0800	tmpfs	tmpfs	/dev/shm
ffff88022bf92640	ffff880233a8a000	devpts	devpts	/dev/pts

Crash工具基本使用

Crash常用命令

ps 查看系统中进程的信息

```
crash> ps
PID      PPID     CPU    TASK                ST  %MEM   VSZ   RSS  COMM
v      0         0      ffffffff81982440  RU   0.0     0     0  [swapper/0]
v      0         0      ffff880233f26810  RU   0.0     0     0  [swapper/1]
v      0         0      ffff880233f273a0  RU   0.0     0     0  [swapper/2]
v      0         0      ffff880233f50000  RU   0.0     0     0  [swapper/3]
v      0         0      ffff880233f50b90  RU   0.0     0     0  [swapper/4]
v      0         0      ffff880233f51720  RU   0.0     0     0  [swapper/5]
v      0         0      ffff880233f522b0  RU   0.0     0     0  [swapper/6]
v      0         0      ffff880233f52e40  RU   0.0     0     0  [swapper/7]
v      0         0      ffff880233f539d0  RU   0.0     0     0  [swapper/8]
v      0         0      ffff880233f54560  RU   0.0     0     0  [swapper/9]
```

mod命令可以显示或者加载内核模块

```
crash> mod
MODULE      NAME                SIZE  OBJECT FILE
ffffffffffa0014680  dm_mod              113547 (not loaded) [CONFIG_KALLSYMS]
ffffffffffa001f1c0  virtio              15008 (not loaded) [CONFIG_KALLSYMS]
ffffffffffa00271e0  serio_raw           13462 (not loaded) [CONFIG_KALLSYMS]
ffffffffffa002d160  dm_log              18411 (not loaded) [CONFIG_KALLSYMS]
ffffffffffa0037080  dm_region_hash      20862 (not loaded) [CONFIG_KALLSYMS]
```

```
crash> mod -s tun /home/linux-rh-3.10/drivers/net/tun.ko
MODULE      NAME      SIZE  OBJECT FILE
ffffffffffa05603c0  tun       27141 /lib/modules/3.10.0-327.77.60.61.x86_64+/kernel/drivers/net/tun.ko
```

search命令可以在内存中寻找所有存放目标数据的地址

```
crash> search kfree_skb
ffff8800017449f8: ffffffff81539250 (kfree_skb)
ffff880001927750: ffffffff81539250 (kfree_skb)
ffff880233df2070: ffffffff81539250 (kfree_skb)
crash>
crash>
crash> search ffff88022b9f8000
ffff88003426c318: ffff88022b9f8000
ffff8800342a3520: ffff88022b9f8000
ffff8800342a3b20: ffff88022b9f8000
```

list命令用于显示链表的内容

```
crash> p file_systems
file_systems = $8 = (struct file_system_type *) 0xffffffff81a0dce0 <sysfs_fs_type>
crash> list file_system_type.next -s file_system_type.name,fs_flags 0xffffffff81a0dce0
0xffffffff81a0dce0
  name = 0xffffffff818c76fb "sysfs"
  fs_flags = 8
0xffffffff81a0e720
  name = 0xffffffff818abc49 "rootfs"
  fs_flags = 0
0xffffffff81a009e0
  name = 0xffffffff818ac0b5 "bdev"
  fs_flags = 0
0xffffffff81a0d760
  name = 0xffffffff8189f7d3 "proc"
  fs_flags = 8
0xffffffff819c1620
  name = 0xffffffff818c17e1 "cgroup"
  fs_flags = 0
0xffffffff819c2c60
  name = 0xffffffff818a6cc6 "cpuset"
  fs_flags = 0
0xffffffff819f31c0
  name = 0xffffffff818d3bb6 "tmpfs"
  fs_flags = 8
```



一个简单例子

查看panic核上的栈

```
crash> bt
PID: 3754  TASK: ffff880034ee8b90  CPU: 3  COMMAND: "systemd-udevd"
#0 [ffff8800b90c3840] machine_kexec at ffffffff81055deb
#1 [ffff8800b90c38a0] crash_kexec at ffffffff810fd992
#2 [ffff8800b90c3970] oops_end at ffffffff81660748
#3 [ffff8800b90c3998] no_context at ffffffff8164feb4
#4 [ffff8800b90c39e8] __bad_area_nosemaphore at ffffffff8164ff4a
#5 [ffff8800b90c3a30] bad_area at ffffffff8165026e
#6 [ffff8800b90c3a58] __do_page_fault at ffffffff816636c9
#7 [ffff8800b90c3ab8] trace_do_page_fault at ffffffff81663816
#8 [ffff8800b90c3af8] do_async_page_fault at ffffffff81662ee9
#9 [ffff8800b90c3b10] async_page_fault at ffffffff8165f6e8
[exception RIP: tun_get_drvinfo+55]
RIP: ffffffff8055c2c7  RSP: ffff8800b90c3bc0  RFLAGS: 00010206
RAX: 0000000000000003  RBX: ffff8800b90c3bdc  RCX: 0000000000000000
RDX: 0000000000000003  RSI: ffffffff8055f0e6  RDI: ffff8800b90c3c03
RBP: ffff8800b90c3bc8  R8: 0000000000000000  R9: ffff3e465a9e7574
R10: 00007fd033b4cc30  R11: 0000000000000246  R12: ffffffff8055f940
R13: ffff8800b90c3bdc  R14: 00007fff3deba7d0  R15: ffffffff81a58fc0
ORIG_RAX: ffffffff81a58fc0  CS: 0010  SS: 0018
#10 [ffff8800b90c3bd0] ethtool_get_drvinfo at ffffffff8154da64
#11 [ffff8800b90c3cd0] dev_ethtool at ffffffff8154f486
#12 [ffff8800b90c3da8] dev_ioctl at ffffffff8155f9ef
#13 [ffff8800b90c3e38] sock_do_ioctl at ffffffff8152c18d
#14 [ffff8800b90c3e60] sock_ioctl at ffffffff8152c8a8
#15 [ffff8800b90c3e90] do_vfs_ioctl at ffffffff81202988
#16 [ffff8800b90c3f00] sys_ioctl at ffffffff81202c11
#17 [ffff8800b90c3f50] system_call_fastpath at ffffffff81668155
RIP: 00007fd033bb5507  RSP: 00007fff3deba788  RFLAGS: 00010246
RAX: 0000000000000010  RBX: 0000556e501678e8  RCX: ffffffff816680a1
RDX: 00007fff3deba7a0  RSI: 00000000000008946  RDI: 0000000000000008
RBP: 00007fff3deba7a0  R8: 0000000000000ffff  R9: 0000000000000007
R10: 00007fd033b4cc30  R11: 0000000000000246  R12: 00007fff3deba7d0
R13: 00007fff3deba970  R14: 0000556e50168600  R15: 0000000000000000
ORIG_RAX: 0000000000000010  CS: 0033  SS: 002b
```

```
crash> log
[43254.289274] BUG: unable to handle kernel NULL pointer dereference at 0000000000000804
[43254.292156] IP: [<ffffffffff8055c2c7>] tun_get_drvinfo+0x37/0x80 [tun]
[43254.293008] PGD 0
[43254.293297] Oops: 0000 [#1] SMP
[43254.293762] Modules linked in: tun(OE) ip6t_rpfilter ipt_REJECT ip6t_REJECT [43254.304526] CPU: 3 PID: 3754 Comm: systemd-udevd Tainted: G
OE ----
 3.10.0-327.77.60.61.x86_64+ #65
[43254.305872] Hardware name: Red Hat KVM, BIOS 0.5.1 01/01/2011
[43254.306599] task: ffff880034ee8b90 ti: ffff8800b90c0000 task.ti: ffff8800b90c0000
[43254.307534] RIP: 0010:[<ffffffffff8055c2c7>] [<ffffffffff8055c2c7>] tun_get_drvinfo+0x37/0x80 [tun]
[43254.308652] RSP: 0018:ffff8800b90c3bc0  EFLAGS: 00010206
[43254.309311] RAX: 0000000000000003  RBX: ffff8800b90c3bdc  RCX: 0000000000000000
[43254.310201] RDX: 0000000000000003  RSI: ffffffff8055f0e6  RDI: ffff8800b90c3c03
[43254.311091] RBP: ffff8800b90c3bc8  R08: 0000000000000000  R09: ffff3e465a9e7574
[43254.311982] R10: 00007fd033b4cc30  R11: 0000000000000246  R12: ffffffff8055f940
[43254.312873] R13: ffff8800b90c3bdc  R14: 00007fff3deba7d0  R15: ffffffff81a58fc0
[43254.313760] FS: 00007fd034f548c0(0000)  GS:ffff8802372c0000(0000)  knlGS:0000000000000000
[43254.314766] CS: 0010  DS: 0000  ES: 0000  CR0: 0000000080050033
[43254.315489] CR2: 0000000000000804  CR3: 00000000357aa000  CR4: 000000000000006e0
[43254.316390] DR0: 0000000000000000  DR1: 0000000000000000  DR2: 0000000000000000
[43254.317271] DR3: 0000000000000000  DR6: 00000000ffff0fff  DR7: 0000000000000400
[43254.318163] Call Trace:
[43254.318513] [<ffffffffff8154da64>] ethtool_get_drvinfo+0x84/0x1d0
[43254.319261] [<ffffffffff8154f486>] dev_ethtool+0xa16/0x1b30
[43254.320015] [<ffffffffff80224dbc>] ? xfs_iunlock+0x11c/0x130 [xfs]
[43254.320796] [<ffffffffff81175efb>] ? unlock_page+0x2b/0x30
[43254.321478] [<ffffffffff815486dc>] ? dev_get_by_name_rcu+0x5c/0x80
[43254.322240] [<ffffffffff8155f9ef>] dev_ioctl+0x1cf/0x590
[43254.322908] [<ffffffffff811a535c>] ? handle_mm_fault+0x65c/0x1010
[43254.323663] [<ffffffffff8152c18d>] sock_do_ioctl+0x4d/0x60
[43254.324333] [<ffffffffff8152c8a8>] sock_ioctl+0x1f8/0x2d0
[43254.325011] [<ffffffffff81202988>] do_vfs_ioctl+0x2e8/0x4d0
[43254.325722] [<ffffffffff81663471>] ? __do_page_fault+0x171/0x430
[43254.326466] [<ffffffffff816680a1>] ? system_call_after_swapgs+0xae/0x146
[43254.327285] [<ffffffffff81202c11>] SyS_ioctl+0xa1/0xc0
[43254.327924] [<ffffffffff81668095>] ? system_call_after_swapgs+0xa2/0x146
[43254.328753] [<ffffffffff81668155>] system_call_fastpath+0x1c/0x21
[43254.329507] [<ffffffffff816680a1>] ? system_call_after_swapgs+0xae/0x146
```

一个简单例子

查看panic核上的栈

```
crash> mod -s tun /home/linux-rh-3.10/drivers/net/tun.ko //加载模块信息
MODULE      NAME      SIZE OBJECT FILE
ffffffa05603c0 tun      27141 /lib/modules/3.10.0-
327.77.60.61.x86_64+/kernel/drivers/net/tun.ko
crash>
crash> dis -l tun_get_drvinfo+0x37 4
/home/linux-rh-3-10/drivers/net/tun.c: 2401
0xffffffffa055c2c7 <tun_get_drvinfo+55>:    mov    0x804,%eax //tun->flags
/home/linux-rh-3-10/drivers/net/tun.c: 2403
0xffffffffa055c2ce <tun_get_drvinfo+62>:    and    $0xf,%eax
0xffffffffa055c2d1 <tun_get_drvinfo+65>:    cmp    $0x1,%eax
0xffffffffa055c2d4 <tun_get_drvinfo+68>:    je     0xffffffffa055c2f8
<tun_get_drvinfo+104>
```

```
2396 static void tun_get_drvinfo(struct net_device *dev, struct ethtool_drvinfo *info)
2397 {
2398     struct tun_struct *tun = netdev_priv(dev);
2399
2400     strlcpy(info->driver, DRV_NAME, sizeof(info->driver));
2401     strlcpy(info->version, DRV_VERSION, sizeof(info->version));
2402
2403     tun = NULL;
2404     switch (tun->flags & TUN_TYPE_MASK) {
2405     case IFF_TUN:
2406         strlcpy(info->bus_info, "tun", sizeof(info->bus_info));
2407         break;
```

```
crash> struct tun_struct -ox
struct tun_struct {
    [0x0] struct tun_file *tfiles[256];
    [0x800] unsigned int numqueues;
    [0x804] unsigned int flags;
    [0x808] kuid_t owner;
    [0x80c] kgid_t group;
    [0x810] struct net_device *dev;
```

```
crash> log
[43254.289274] BUG: unable to handle kernel NULL pointer dereference at 0000000000000804
[43254.292156] IP: [<ffffffa055c2c7>] tun_get_drvinfo+0x37/0x80 [tun]
[43254.293008] PGD 0
[43254.293297] Oops: 0000 [#1] SMP
[43254.293762] Modules linked in: tun(OE) ip6t_rpfilter ipt_REJECT ip6t_REJECT [43254.304526] CPU: 3 PID: 3754 Comm: systemd-udevd Tainted: G
OE ----
3.10.0-327.77.60.61.x86_64+ #65
[43254.305872] Hardware name: Red Hat KVM, BIOS 0.5.1 01/01/2011
[43254.306599] task: ffff880034ee8b90 ti: ffff8800b90c0000 task.ti: ffff8800b90c0000
[43254.307534] RIP: 0010:[<ffffffa055c2c7>] [<ffffffa055c2c7>] tun_get_drvinfo+0x37/0x80 [tun]
[43254.308652] RSP: 0018:ffff8800b90c3bc0 EFLAGS: 00010206
[43254.309311] RAX: 0000000000000003 RBX: ffff8800b90c3bdc RCX: 0000000000000000
[43254.310201] RDX: 0000000000000003 RSI: ffffffa055f0e6 RDI: ffff8800b90c3c03
[43254.311091] RBP: ffff8800b90c3bc8 R08: 0000000000000000 R09: ffff3e465a9e7574
[43254.311982] R10: 00007fd033b4cc30 R11: 0000000000000246 R12: ffffffa055f940
[43254.312873] R13: ffff8800b90c3bdc R14: 00007fff3deba7d0 R15: ffffff81a58fc0
[43254.313760] FS: 00007fd034f548c0(0000) GS:ffff8802372c0000(0000) knlGS:0000000000000000
[43254.314766] CS: 0010 DS: 0000 ES: 0000 CR0: 0000000080050033
[43254.315489] CR2: 0000000000000804 CR3: 00000000357aa000 CR4: 00000000000006e0
[43254.316390] DR0: 0000000000000000 DR1: 0000000000000000 DR2: 0000000000000000
[43254.317271] DR3: 0000000000000000 DR6: 00000000ffff0ff0 DR7: 0000000000000400
[43254.318163] Call Trace:
[43254.318513] [<fffffff8154da64>] ethtool_get_drvinfo+0x84/0x1d0
[43254.319261] [<fffffff8154f486>] dev_ethtool+0xa16/0x1b30
[43254.320015] [<ffffffa0224dbc>] ? xfs_iunlock+0x11c/0x130 [xfs]
[43254.320796] [<fffffff81175efb>] ? unlock_page+0x2b/0x30
[43254.321478] [<fffffff815486dc>] ? dev_get_by_name_rcu+0x5c/0x80
[43254.322240] [<fffffff8155f9ef>] dev_ioctl+0x1cf/0x590
[43254.322908] [<fffffff811a535c>] ? handle_mm_fault+0x65c/0x1010
[43254.323663] [<fffffff8152c18d>] sock_do_ioctl+0x4d/0x60
[43254.324333] [<fffffff8152c8a8>] sock_ioctl+0x1f8/0x2d0
[43254.325011] [<fffffff81202988>] do_vfs_ioctl+0x2e8/0x4d0
[43254.325722] [<fffffff81663471>] ? __do_page_fault+0x171/0x430
[43254.326466] [<fffffff816680a1>] ? system_call_after_swapgs+0xae/0x146
[43254.327285] [<fffffff81202c11>] SyS_ioctl+0xa1/0xc0
[43254.327924] [<fffffff81668095>] ? system_call_after_swapgs+0xa2/0x146
[43254.328753] [<fffffff81668155>] system_call_fastpath+0x1c/0x21
[43254.329507] [<fffffff816680a1>] ? system_call_after_swapgs+0xae/0x146
```



现网问题实战

查看panic核上的栈，然后反汇编找出问题点

```
crash> bt
PID: 404708 TASK: ffff803de8094040 CPU: 27 COMMAND: "java"
#0 [ffff803ffff3e950] crash_kexec at ffff000008176a48
#1 [ffff803ffff3e980] die at ffff00000808d8fc
#2 [ffff803ffff3e9c0] die_kernel_fault at ffff0000080a4540
#3 [ffff803ffff3e9f0] __do_kernel_fault at ffff0000080a42b4
#4 [ffff803ffff3ea20] do_page_fault at ffff0000089ffc60
#5 [ffff803ffff3eb10] do_translation_fault at ffff000008a00158
#6 [ffff803ffff3eb20] do_mem_abort at ffff000008081260
#7 [ffff803ffff3ed00] ell_ia at ffff000008083114
    PC: ffff000008938550 [inet_sock_destruct+64]
    LR: ffff00000887b83c [__sk_destruct+44]
    SP: ffff803ffff3ed10 PSTATE: 80400009
    X29: ffff803ffff3ed10 X28: ffff803f26209b00 X27: ffff803f263c0400
    X26: ffff803f26209840 X25: 0000000000002317 X24: ffff803f26209900
    X23: ffff000001991d96 X22: 0000000000002441 X21: ffff803f7489dee8
    X20: ffff803f7489e0c8 X19: ffff803f7489e000 X18: 0000000000000000
    X17: 0000000000000000 X16: ffffffff803fc4b8 X15: ffffffff803fc400
    X14: ffffffff803fddb8 X13: 0000000000000000 X12: 0000000000000000
    X11: 0000000000000000 X10: 0000000000000040 X9: ffff0000080a26e0
    X8: 0000000000000006 X7: 0000000000000000 X6: 0000000000000002
    X5: 0000000000000000 X4: 0000000000000020 X3: ffff805fc86fb010
    X2: ffff803fc4817000 X1: 0000000000000000 X0: ffff803fc481ab00
#8 [ffff803ffff3ed10] inet_sock_destruct at ffff00000893854c
#9 [ffff803ffff3ed30] __sk_destruct at ffff00000887b838
#10 [ffff803ffff3ed60] sk_destruct at ffff00000887c624
#11 [ffff803ffff3ed80] __sk_free at ffff00000887c674
#12 [ffff803ffff3eda0] __sock_wfree at ffff00000887ce94
#13 [ffff803ffff3edc0] skb_release_head_state at ffff000008883128
#14 [ffff803ffff3ee10] skb_release_all at ffff000008883200
#15 [ffff803ffff3ee30] napi_consume_skb at ffff000008886de4
#16 [ffff803ffff3ee50] mlx5e_poll_tx_cq at ffff000001950a44 [mlx5_core]
#17 [ffff803ffff3eef0] mlx5e_napi_poll at ffff000001955264 [mlx5_core]
#18 [ffff803ffff3ef90] net_rx_action at ffff00000889c3d0
#19 [ffff803ffff3f020] __softirqentry_text_start at ffff0000080818c0
#20 [ffff803ffff3f0b0] irq_exit at ffff0000080dc680
#21 [ffff803ffff3f0d0] __handle_domain_irq at ffff000008136a94
#22 [ffff803ffff3f110] gic_handle_irq at ffff000008081678
```

现网问题实战

反汇编出问题的函数inet_sock_destruct, 分析代码流程

```
/home/jiangdi/luyun/klinux-4.19/net/ipv4/af_inet.c: 135
0xffff00008938510 <inet_sock_destruct>:      stp     x29, x30, [sp,#-32]!
0xffff00008938514 <inet_sock_destruct+4>:      mov     x29, sp
0xffff00008938518 <inet_sock_destruct+8>:      stp     x19, x20, [sp,#16]
/home/jiangdi/luyun/klinux-4.19/net/ipv4/af_inet.c: 138
0xffff0000893851c <inet_sock_destruct+12>:      add     x20, x0, #0xc8      x0是sock, x20放receive队列head
/home/jiangdi/luyun/klinux-4.19/net/ipv4/af_inet.c: 135
0xffff00008938520 <inet_sock_destruct+16>:      mov     x19, x0      x19保存sock
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 1631
0xffff00008938524 <inet_sock_destruct+20>:      ldr     x0, [x0,#200]      x0放置receive队列next
                             指向的skb
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 1633
0xffff00008938528 <inet_sock_destruct+24>:      cmp     x20, x0
0xffff0000893852c <inet_sock_destruct+28>:      b.eq    0xffff00008938564 <inet_sock_destruct+84>
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 1917
0xffff00008938530 <inet_sock_destruct+32>:      cbz     x0, 0xffff00008938564 <inet_sock_destruct+84>
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 1896
0xffff00008938534 <inet_sock_destruct+36>:      ldr     w1, [x20,#16]
0xffff00008938538 <inet_sock_destruct+40>:      sub     w1, w1, #0x1      len--
/home/jiangdi/luyun/klinux-4.19/.include/linux/compiler.h: 219
0xffff0000893853c <inet_sock_destruct+44>:      str     w1, [x19,#216]
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 1898
0xffff00008938540 <inet_sock_destruct+48>:      ldp     x2, x1, [x0]      x1放入next, x2放入prev
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 1899
0xffff00008938544 <inet_sock_destruct+52>:      dmb     ish
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 1900
0xffff00008938548 <inet_sock_destruct+56>:      stp     xzr, xzr, [x0]
/home/jiangdi/luyun/klinux-4.19/.include/linux/compiler.h: 220
0xffff0000893854c <inet_sock_destruct+60>:      str     x1, [x2,#8]      x1放入x2指向的skb的prev
                             指针
0xffff00008938550 <inet_sock_destruct+64>:      str     x2, [x1]          访问x1, 发空指针
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 2636
0xffff00008938554 <inet_sock_destruct+68>:      bl      0xffff00008883248 <kfree_skb>
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 1631
0xffff00008938558 <inet_sock_destruct+72>:      ldr     x0, [x20]
/home/jiangdi/luyun/klinux-4.19/.include/linux/skbuff.h: 1633
0xffff0000893855c <inet_sock_destruct+76>:      cmp     x20, x0
0xffff00008938560 <inet_sock_destruct+80>:      b.ne    0xffff00008938530 <inet_sock_destruct+32>
```

现网问题实战

可以看到问题是在访问x2寄存器保存的skb地址
偏移8地址时发生了空指针引用

```
crash> sk_buff ffff803fc4817000 -o
struct sk_buff {
    union {
        struct {
            [ffff803fc4817000] struct sk_buff *next;
            [ffff803fc4817008] struct sk_buff *prev;
            union {
                [ffff803fc4817010] struct net_device *dev;
                [ffff803fc4817010] unsigned long dev_scratch;
            };
        };
        [ffff803fc4817000] struct rb_node rbnode;
        [ffff803fc4817000] struct list_head list;
    };
};
```

```
crash> sk_buff ffff803fc4817000
struct sk_buff {
    {
        {
            next = 0xffff803fc4818600,
            prev = 0x0,
            {
                dev = 0xffff803f9c940000,
                dev_scratch = 18446603609431080960
            }
        },
        rbnode = {
            __rb_parent_color = 18446603610100958720,
            rb_right = 0x0,
            rb_left = 0xffff803f9c940000
        },
        list = {
            next = 0xffff803fc4818600,
            prev = 0x0
        }
    },
    {
        sk = 0xffff803f098b7c00,
        ip_defrag_offset = 160136192
    },
};
```


现网问题实战

查看对应的C代码

```
void inet_sock_destruct(struct sock *sk)
{
    struct inet_sock *inet = inet_sk(sk);

    __skb_queue_purge(&sk->sk_receive_queue);
    __skb_queue_purge(&sk->sk_error_queue);

    sk_mem_reclaim(sk);

    if (sk->sk_type == SOCK_STREAM && sk->sk_state != TCP_CLOSE) {
        pr_err("Attempt to release TCP socket in state %d %p\n",
            ... sk->sk_state, sk);
        return;
    }
    if (!sock_flag(sk, SOCK_DEAD)) {
        pr_err("Attempt to release alive inet socket %p\n", sk);
        return;
    }
}

WARN_ON(atomic_read(&sk->sk_rmem_alloc));
WARN_ON(refcount_read(&sk->sk_wmem_alloc));
WARN_ON(sk->sk_wmem_queued);
WARN_ON(sk->sk_forward_alloc);

kfree(rcu_dereference_protected(inet->inet_opt, 1));
dst_release(rcu_dereference_protected(sk->sk_dst_cache, 1));
dst_release(sk->sk_rx_dst);
sk_refcnt_debug_dec(sk);
}
EXPORT_SYMBOL(inet_sock_destruct);
```

```
static inline void __skb_queue_purge(struct sk_buff_head *list)
{
    struct sk_buff *skb;
    while ((skb = __skb_dequeue(list)) != NULL)
        kfree_skb(skb);
}
```

```
static inline void __skb_unlink(struct sk_buff *skb, struct sk_buff_head *list)
{
    struct sk_buff *next, *prev;

    WRITE_ONCE(list->qlen, list->qlen - 1);
    next = skb->next;
    prev = skb->prev;
    skb->next = skb->prev = NULL;
    WRITE_ONCE(next->prev, prev);
    WRITE_ONCE(prev->next, next);
}
```

现网问题实战

从上面反汇编代码可以知道，X20寄存器存放的是sk->sk_receive_queue，解析可以看到qlen为0，说明receive队列的报文已经释放过

再查看x2寄存器保存的skb发现是其他连接的报文，与当前释放的sock不同，所以怀疑原来的报文已经被释放过一次

```
crash> sk_buff_head ffff803f7489e0c8
struct sk_buff_head {
  next = 0xffff803f7489e0c8,
  prev = 0xffff803f7489e0c8,
  qlen = 0,
  lock = {
    {
      rlock = {
        raw_lock = {
          {
            val = {
              counter = 0
            },
            {
              locked = 0 '\000',
              pending = 0 '\000'
            },
            {
              locked_pending = 0,
              tail = 0
            }
          }
        }
      }
    }
  }
}
```

现网问题实战

打上一个调测补丁，在skb报文释放时保存栈到skb->cb区域，然后尝试复现

```
--- net/core/skbuff.c.orig 2020-11-21 11:01:44.569912940 +0800
+++ net/core/skbuff.c 2020-11-21 11:14:20.009912940 +0800
@@ -607,6 +607,17 @@ fastpath:
     kmem_cache_free(skbuff_fclone_cache, fclones);
 }

+static void __save_stack_trace(unsigned long *trace)
+{
+    struct stack_trace stack_trace;
+
+    stack_trace.max_entries = 5;
+    stack_trace.nr_entries = 0;
+    stack_trace.entries = trace;
+    stack_trace.skip = 1;
+    save_stack_trace(&stack_trace);
+}
+
void skb_release_head_state(struct sk_buff *skb)
{
    skb_dst_drop(skb);
@@ -615,6 +626,7 @@ void skb_release_head_state(struct sk_buff
    WARN_ON(in_irq());
    skb->destructor(skb);
}
+    __save_stack_trace((unsigned long *)skb->cb);
#ifdef CONFIG_NF_CONNTRACK
    nf_conntrack_put(skb_nfct(skb));
#endif
```


现网问题实战

问题出现时读取receive队列的skb(即x0)的cb区域,

从中可以看到保存的报文之前释放的路径:

tcp_done -> inet_csk_destroy_sock->skb_release_all->skb_release_all->kfree_skb

```
crash> sk_buff.cb -o ffff803fc481ab00
struct sk_buff {
  [ffff803fc481ab28] char cb[48];
}
crash> rd -s ffff803fc481ab28 20
ffff803fc481ab28:  skb_release_all+20 kfree_skb+44
ffff803fc481ab38:  sk_stream_kill_queues+72 inet_csk_destroy_sock+88
ffff803fc481ab48:  tcp_done+280      0000000000000000
ffff803fc481ab58:  0000000000000000 sock_rfree
ffff803fc481ab68:  0000000000000000 0000000000000000
ffff803fc481ab78:  0000000000000000 0000000000000000
ffff803fc481ab88:  002000010000000e 00000000000008140
ffff803fc481ab98:  000000000000092a5 9e9e9e9e0000000a
ffff803fc481aba8:  000004f500000000 0000000000000000
ffff803fc481abb8:  0000000000000000 0042005000640008
```

现网问题实战

对照vmcore的栈

问题也发生在对x2寄存器保存skb的访问

```
#4 [ffff803fffc7e6c0] do_page_fault at ffff0000089ffc60
#5 [ffff803fffc7e7b0] do_translation_fault at ffff000008a00158
#6 [ffff803fffc7e7c0] do_mem_abort at ffff000008081260
#7 [ffff803fffc7e9a0] ell_ia at ffff000008083114
    PC: ffff00000888a610 [sk_stream_kill_queues+64]
    LR: ffff00000888a618 [sk_stream_kill_queues+72]
    SP: ffff803fffc7e9b0 PSTATE: 20400009
    X29: ffff803fffc7e9b0 X28: 0000000000000001 X27: 0000000000000042
    X26: ffff803fc7346051 X25: 0000000000000000 X24: ffff803fc7346064
    X23: 0000000000000006 X22: 0000000000000000 X21: ffff803fcb9c94728
    X20: fffffa03fd978f0c8 X19: fffffa03fd978f000 X18: 0000000000000001
    X17: 0000000000000000 X16: ffffffff803f4b8 X15: ffffffff803f4b8
    X14: ffff000008935940f X13: ffff0000089359422 X12: ffff0000091de000
    X11: ffff0000091c57a0 X10: ffff000008629738 X9: 0000000000000000
    X8: 00000000000005dfa X7: 0000000000000000 X6: 0000000000018d44
    X5: 0000000000000000 X4: fffffa03fd978f0c8 X3: 0000000000018d43
    X2: fffffa03fd978f0c8 X1: 0000000000000000 X0: ffff803fcb9c94700
#8 [ffff803fffc7e9b0] sk_stream_kill_queues at ffff00000888a60c
#9 [ffff803fffc7e9d0] inet_csk_destroy_sock at ffff0000089026ac
#10 [ffff803fffc7e9f0] tcp_done at ffff0000089054bc
#11 [ffff803fffc7ea10] tcp_time_wait at ffff000008921518
#12 [ffff803fffc7ea50] tcp_fin at ffff0000089120d0
#13 [ffff803fffc7ea70] tcp_data_queue at ffff00000891367c
#14 [ffff803fffc7eaf0] tcp_rcv_state_process at ffff0000089141f4
#15 [ffff803fffc7eb60] tcp_v4_do_rcv at ffff00000891dbe8
#16 [ffff803fffc7ebb0] tcp_v4_rcv at ffff000008921274
#17 [ffff803fffc7ec40] ip_local_deliver_finish at ffff0000088f5c48
#18 [ffff803fffc7ec80] ip_local_deliver at ffff0000088f6960
#19 [ffff803fffc7ecd0] ip_rcv_finish at ffff0000088f6254
#20 [ffff803fffc7ed00] ip_rcv at ffff0000088f6a30
#21 [ffff803fffc7ed60] __netif_receive_skb_one_core at ffff00000889bdd8
#22 [ffff803fffc7eda0] __netif_receive_skb at ffff00000889be24
#23 [ffff803fffc7edc0] netif_receive_skb_internal at ffff00000889bec0
#24 [ffff803fffc7edf0] napi_gro_receive at ffff00000889cc0c
#25 [ffff803fffc7ee10] mlx5e_handle_rx_cqe at ffff0000019531dc [mlx5_core]
#26 [ffff803fffc7ee80] mlx5e_poll_rx_cq at ffff0000019548a0 [mlx5_core]
```

```
crash> dis -l sk_stream_kill_queues
/home/jiangdi/luyun/klinux-4.19/net/core/stream.c: 194
0xffff00000888a5d0 <sk_stream_kill_queues>: stp x29, x30, [sp,#-32]!
0xffff00000888a5d4 <sk_stream_kill_queues+4>: mov x29, sp
0xffff00000888a5d8 <sk_stream_kill_queues+8>: stp x19, x20, [sp,#16]
/home/jiangdi/luyun/klinux-4.19/net/core/stream.c: 196
0xffff00000888a5dc <sk_stream_kill_queues+12>: add x20, x0, #0xc8
/home/jiangdi/luyun/klinux-4.19/net/core/stream.c: 194
0xffff00000888a5e0 <sk_stream_kill_queues+16>: mov x19, x0
/home/jiangdi/luyun/klinux-4.19/include/linux/skbuff.h: 1631
0xffff00000888a5e4 <sk_stream_kill_queues+20>: ldr x0, [x0,#200]
/home/jiangdi/luyun/klinux-4.19/include/linux/skbuff.h: 1633
0xffff00000888a5e8 <sk_stream_kill_queues+24>: cmp x20, x0
0xffff00000888a5ec <sk_stream_kill_queues+28>: b.eq 0xffff00000888a624 <sk_stream_k
/home/jiangdi/luyun/klinux-4.19/include/linux/skbuff.h: 1917
0xffff00000888a5f0 <sk_stream_kill_queues+32>: cbz x0, 0xffff00000888a624 <sk_stre
/home/jiangdi/luyun/klinux-4.19/include/linux/skbuff.h: 1896
0xffff00000888a5f4 <sk_stream_kill_queues+36>: ldr w1, [x20,#16]
0xffff00000888a5f8 <sk_stream_kill_queues+40>: sub w1, w1, #0x1
/home/jiangdi/luyun/klinux-4.19/include/linux/compiler.h: 219
0xffff00000888a5fc <sk_stream_kill_queues+44>: str w1, [x19,#216]
/home/jiangdi/luyun/klinux-4.19/include/linux/skbuff.h: 1898
0xffff00000888a600 <sk_stream_kill_queues+48>: ldp x2, x1, [x0]
/home/jiangdi/luyun/klinux-4.19/include/linux/skbuff.h: 1899
0xffff00000888a604 <sk_stream_kill_queues+52>: dmb ish
/home/jiangdi/luyun/klinux-4.19/include/linux/skbuff.h: 1900
0xffff00000888a608 <sk_stream_kill_queues+56>: stp xzr, xzr, [x0]
/home/jiangdi/luyun/klinux-4.19/include/linux/compiler.h: 220
0xffff00000888a60c <sk_stream_kill_queues+60>: str x1, [x2,#8]
0xffff00000888a610 <sk_stream_kill_queues+64>: str x2, [x1]
/home/jiangdi/luyun/klinux-4.19/include/linux/skbuff.h: 2636
0xffff00000888a614 <sk_stream_kill_queues+68>: bl 0xffff000008883248 <kfree_skb>
/home/jiangdi/luyun/klinux-4.19/include/linux/skbuff.h: 1631
0xffff00000888a618 <sk_stream_kill_queues+72>: ldr x0, [x20]
```

现网问题实战

解析x2寄存器中的skb报文，可以看到prev指针也为空

```
crash> sk_buff ffffa03fd978f0c8
struct sk_buff {
{
{
next = 0xffff803fcbc94700,
prev = 0x0,
{
dev = 0xffffffff,
dev_scratch = 4294967295
}
},
rbnode = {
__rb_parent_color = 18446603610223101696,
rb_right = 0x0,
rb_left = 0xffffffff
},
list = {
next = 0xffff803fcbc94700,
prev = 0x0
}
},
{
sk = 0x300,
ip_defrag_offset = 768
},
}
```


现网问题实战

从上述两个vmcore的栈分析，可以看到存在
两条并发路径操作同一sock的receive队列：

```
cpu0
-----
net_rx_action
-->mlx5e_napi_poll
-->mlx5e_poll_rx_cq
--->mlx5_handle_rx_cqe
--->napi_gro_receive
--->netif_receive_skb_internal
-->__netif_receive_skb
-->__netif_receive_skb_one_core
-->ip_rcv
-->ip_rcv_finish
-->ip_local_deliver
-->ip_local_deliver_finish
-->tcp_v4_rcv
-->tcp_v4_do_rcv
-->tcp_rcv_state_process
-->tcp_data_queue
-->tcp_fin //收到FIN
-->tcp_time_wait //发送ACK之后迁移状态
-->tcp_done //TCP_CLOSE
-->inet_csk_destroy_sock
-->sk_stream_kill_queues
-->__skb_queue_purge(&sk->sk_receive_queue) //释放报文
-->sock_put
-->sk_free
-->if (refcount_dec_and_test(&sk->sk_wmem_alloc))
```

```
cpu1
-----
net_rx_action
-->mlx5e_napi_poll
-->mlx5e_poll_tx_cq
--->napi_consume_skb
-->skb_release_all
-->skb_release_head_state
-->skb->destructor //此处destructor 为__sock_wfree
//cpu0流程同时并发执行完
-->__sock_wfree
-->if (refcount_sub_and_test(skb->truesize, &sk->sk_wmem_alloc))
__sk_free
-->sk_destruct
-->__sk_destruct
-->inet_sock_destruct
-->__skb_queue_purge(&sk->sk_receive_queue)
//从第一个栈的分析中知道，此时receive queue已经在cpu0上被清空
//但此处cpu1还是从sock获取到receive 队列里的报文再次去释放而发生问题
//所以分析问题原因cpu0释放skb后未刷入内存导致cpu1获取到旧数据
```

现网问题实战

根因总结：

网卡在对同一条tcp流做hash时收发分发到了不同的CPU上处理，
cpu1在收到FIN报文处理流程中会调用__skb_queue_purge清空释放sock的receive队列，
同时cpu0上在释放回收发送的ACK报文时也会调用__skb_queue_purge清空receive队列。
但在cpu0 __sock_wfree之前没有内存屏障，存在内存序问题，
导致cpu0获取到sk_receive_queued的过期数据，从而触发double free。

抽象简单模型如下：

cpu0	cpu1

	<u>bh_lock_sock_nested(sk);</u>
	<u>__skb_queue_purge(&sk->sk_receive_queue);</u> //A

	<u>if (refcount_dec_and_test(&sk->sk_wmem_alloc))</u> //B

	<u>bh_unlock_sock(sk)</u> //spinlock 锁,已包含 barrier 语义

<u>__sock_wfree</u>	
<u>if (refcount_sub_and_test(skb->truesize, &sk->sk_wmem_alloc))</u> //B	
<u>__sk_free</u>	
<u>sk_destruct</u>	
<u>__sk_destruct</u>	
<u>inet_sock_destruct</u>	
<u>__skb_queue_purge(&sk->sk_receive_queue);</u> //A	

上述模型中，如果 cpu0 上的流程先于 cpu1 执行，则不会发生问题。因为在 cpu0 上 sk_wmem_alloc 没有减到 0，不会触发 __sk_free，receive 队列仅在 cpu1 上释放一次。

解决方案：

__sock_wfree在执行sk_free前需要加内存屏障

Thank you