



# Linux PCI子系统介绍

汪雄峰

# 目录

CONTENT

**01** 硬件基础

---

**02** 基本概念介绍

---

**03** 配置空间介绍

---

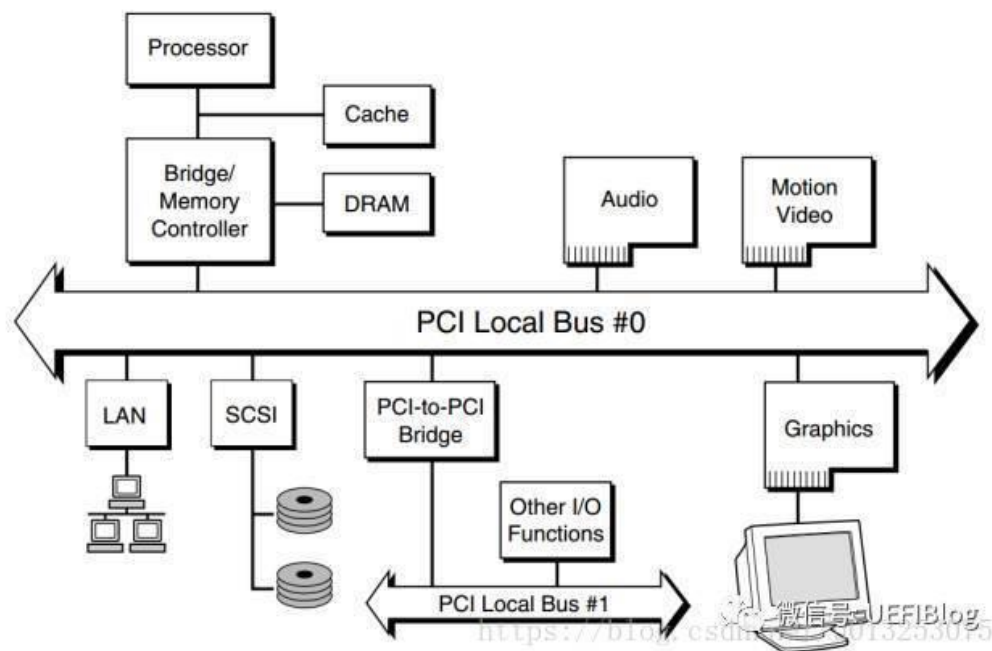

**04** 初始化流程

---

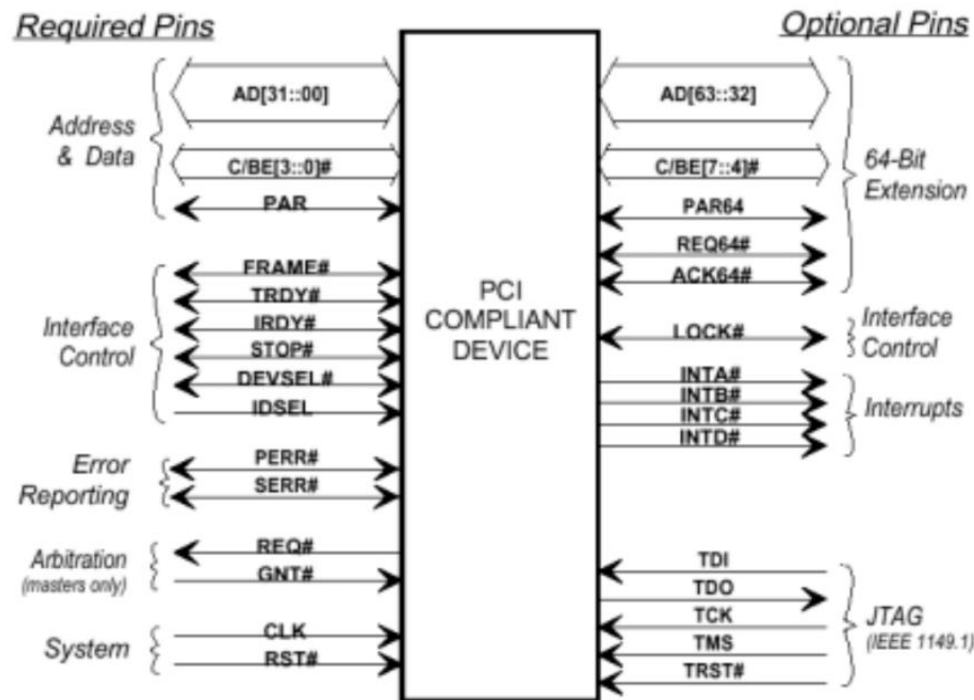


# 什么是PCI设备

- PCI总线
- 并行传输
- 仲裁机制
- 共享总线结构



一颗典型的PCI总线树





PCI总线引脚图


# PCI与PCIe

## PCIe总线:

- 差分串行传输
- 点对点结构
- 多条lane



引脚号	定义 (B)	说明	定义 (A)	说明
1	+12V	+12V电压	PRSTN1#	热拔插存在检测
2	+12V	+12V电压	+12V	+12V电压
3	RSVD	保留引脚	+12V	+12V电压
4	GND	地	GND	地
5	SMCLK	系统管理总线时钟	JTAG2	测试时钟/TCK
6	SMDAT	系统管理总线数据	JTAG3	测试数据输入/TDI
7	GND	地	JTAG4	测试数据输出/TDO
8	+3.3V	+3.3V电压	JTAG5	测试模式选择/TMS
9	JTAG1	测试复位/TRST	+3.3V	+3.3V电压
10	3.3VAUX	3.3V辅助电源	+3.3V	+3.3V电压
11	WAKE#	链接激活信号	PWRGD	电源准备好信号
12	RSVD	保留引脚	GND	地
13	GND	地	REFCLK+	差分信号对参考时钟
14	HSOp(0)	0号信道发送差分信号对	REFCLK-	地
15	HSOn(0)		GND	地
16	GND	地	HSIp(0)	0号信道接收差分信号对
17	PRSTN2#	热拔插存在检测	HSIn(0)	地
18	GND	地	GND	地




PCIe总线引脚图

# PCI与PCIe

## PCIe总线:

- 差分串行传输
- 点对点结构
- 多条lane



针脚号	定义 (B)
1	+12V
2	+12V
3	RSVD
4	GND
5	SMCLK
6	SMDAT
7	GND
8	+3.3V
9	JTAG1
10	3.3VAUX
11	WAKE#
12	RSVD
13	GND
14	HS0p(0)
15	HS0n(0)
16	GND
17	PRSNT2#
18	GND

标准	时钟	传输位宽	每时钟数据	带宽
ISA	4.77 MHz	8	1	4.77 MB/s
ISA	8 MHz	16	0.5	8 MB/s
MCA	5 MHz	16	1	10 MB/s
MCA	5 MHz	32	1	20 MB/s
EISA	8.33 MHz	32	1	33.3 MB/s (16.7 MB/s typically)
VLB	33 MHz	32	1	133 MB/s
PCI	33 MHz	32	1	133 MB/s
PCI-X 66	66 MHz	64	1	533 MB/s
PCI-X 133	133 MHz	64	1	1,066 MB/s
PCI-X 266	133 MHz	64	2	2,132 MB/s
PCI-X 533	133 MHz	64	4	4,266 MB/s
AGP x1	66 MHz	32	1	266 MB/s
AGP x2	66 MHz	32	2	533 MB/s
AGP x4	66 MHz	32	4	1,066 MB/s
AGP x8	66 MHz	32	8	2,133 MB/s
PCIe 1.0 x1	2.5 GHz	1	1	250 MB/s
PCIe 1.0 x4	2.5 GHz	4	1	1,000 MB/s
PCIe 1.0 x8	2.5 GHz	8	1	2,000 MB/s
PCIe 1.0 x16	2.5 GHz	16	1	4,000 MB/s
PCIe 2.0 x1	5 GHz	1	1	500 MB/s
PCIe 2.0 x4	5 GHz	4	1	2,000 MB/s
PCIe 2.0 x8	5 GHz	8	1	4,000 MB/s
PCIe 2.0 x16	5 GHz	16	1	8,000 MB/s
PCIe 3.0 x1	8 GHz	1	1	1,000 MB/s
PCIe 3.0 x4	8 GHz	4	1	4,000 MB/s
PCIe 3.0 x8	8 GHz	8	1	8,000 MB/s
PCIe 3.0 x16	8 GHz	16	1	16,000 MB/s

PCI与PCIe总线速度对比






# PCIe硬件形态



SSD

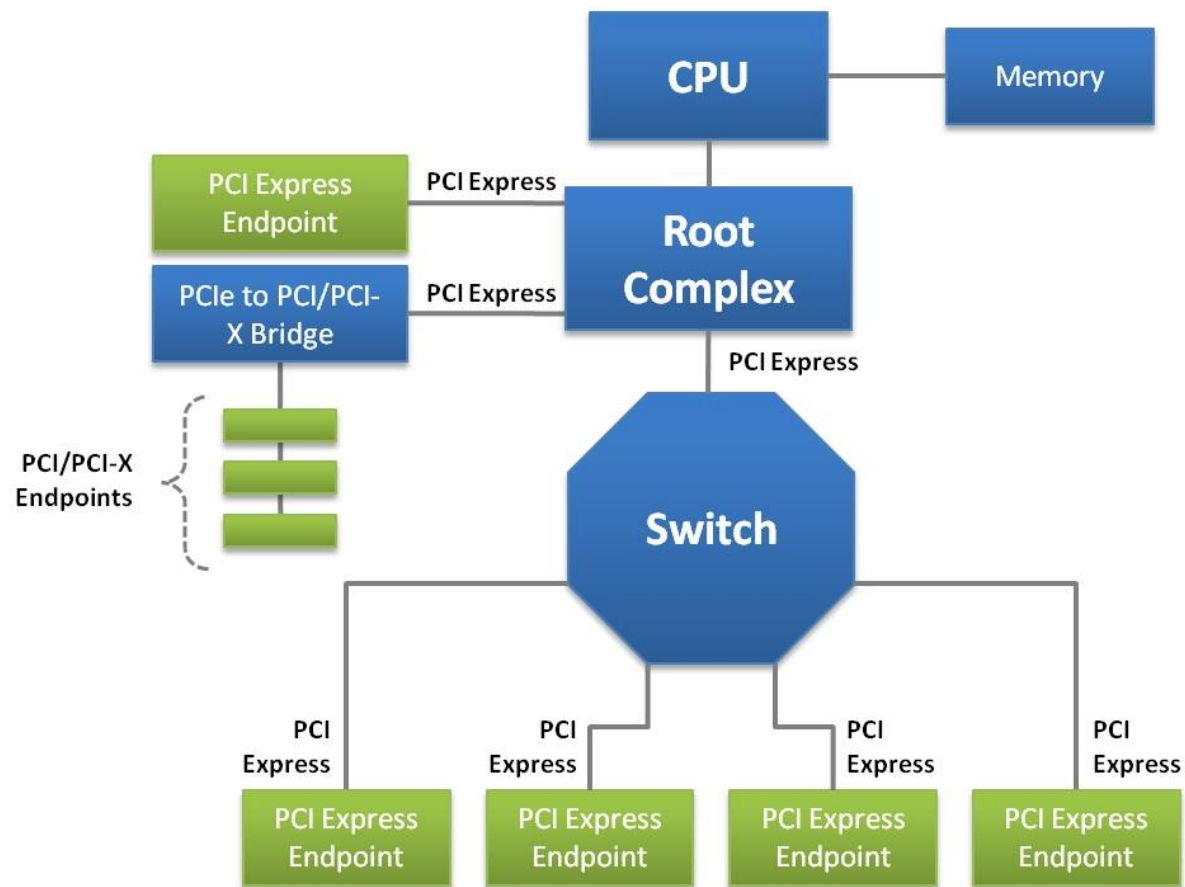


网卡



PCI与PCIe插槽

# PCIe总线拓扑

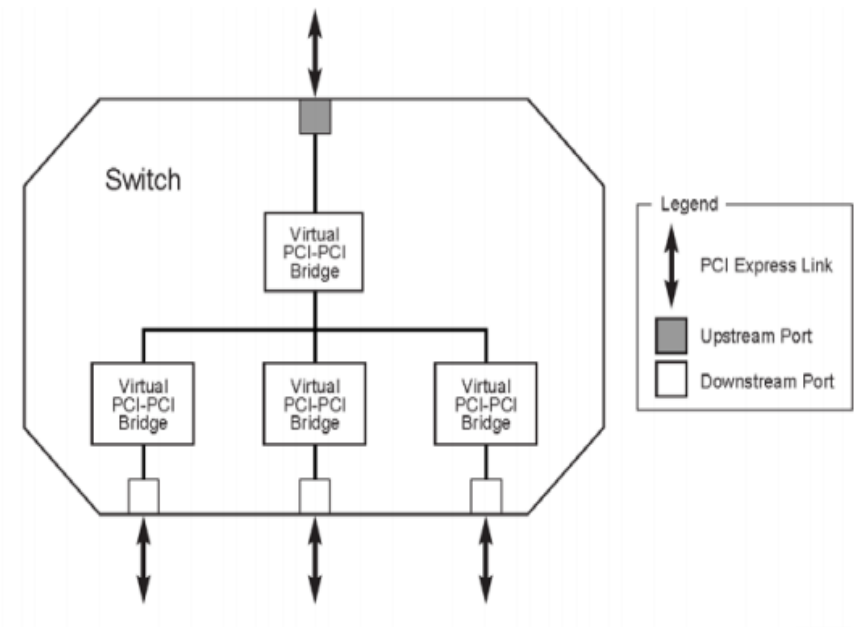
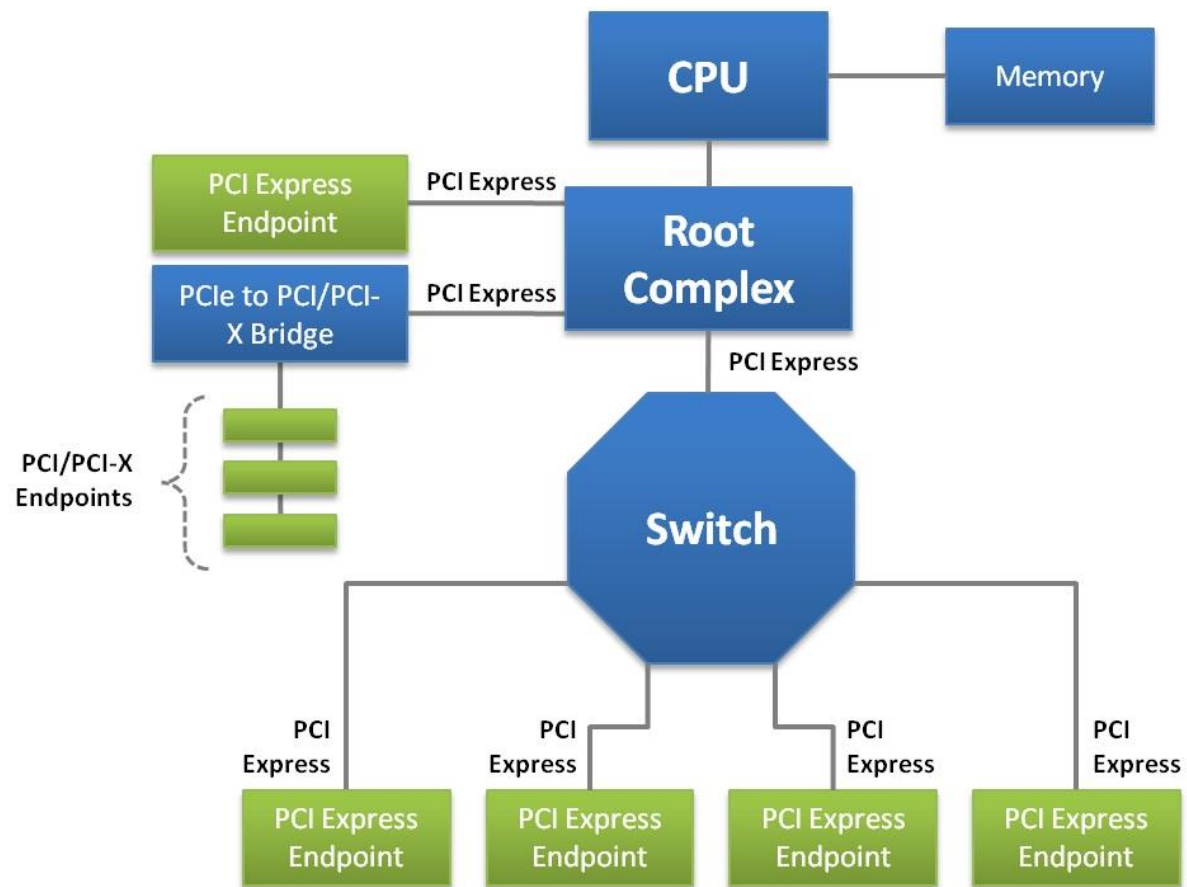


- Root Complex (RC)
- Switch
- Root Port (RP)
- Upstream Port
- Downstream Port
- Endpoint (EP)

为什么需要Switch ?

Switch内部结构 ?

# PCIe总线拓扑



为什么需要Switch？

Switch内部结构？



# CPU域与PCI域

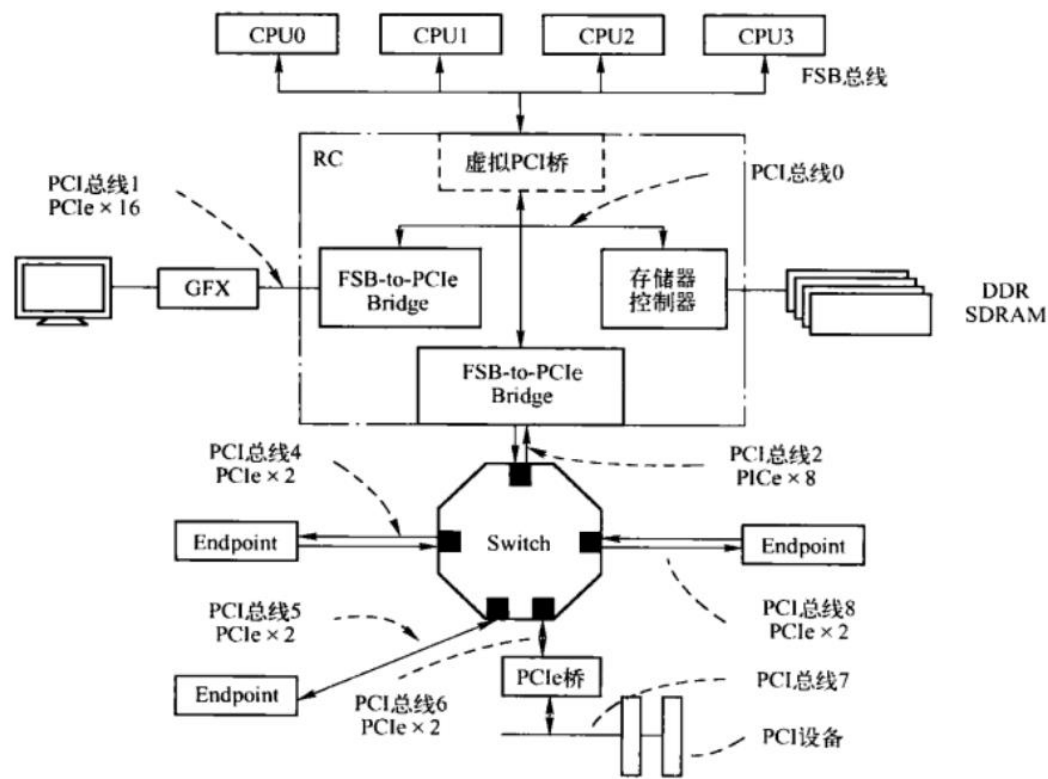



图 4-7 基于 PCIe 总线的处理器系统 A



- 事务层
- 数据链路层
- 物理层

TLP Prefix (Optional)	TLP Prefix (Optional)	TLP Head	Data Payload	TLP Digest(Optional)
-----------------------	-----------------------	----------	--------------	----------------------

图 6-1 TLP 的格式

Switch怎么判断数据报文发到下行端口中的哪一个？

# 查看系统PCI设备

- 查看pci设备拓扑: `lspci -tv (-Dnn)`
- 查看配置空间: `lspci -vvv -s BDF号`
- 以二进制方式查看配置空间: `lspci -xxx -s BDF号 (-xxxx 查看扩展配置空间)`

```
[root@server1 ~]# lspci -tv -Dnn
+-[0000:7c]---00.0-[7d]---00.0  Huawei Technologies Co., Ltd. HNS GE/10GE/25GE RDMA Network Controller [19e5:a222]
|
|      +-00.1  Huawei Technologies Co., Ltd. HNS GE/10GE/25GE Network Controller [19e5:a221]
|
|      +-00.2  Huawei Technologies Co., Ltd. HNS GE/10GE/25GE RDMA Network Controller [19e5:a222]
|
|      \-00.3  Huawei Technologies Co., Ltd. HNS GE/10GE/25GE Network Controller [19e5:a221]
+-[0000:7b]---00.0  Huawei Technologies Co., Ltd. HiSilicon Embedded DMA Engine [19e5:a122]
+-[0000:7a]---00.0  Huawei Technologies Co., Ltd. HiSilicon USB 1.1 Host Controller [19e5:a23b]
|
|      +-01.0  Huawei Technologies Co., Ltd. HiSilicon USB 2.0 2-port Host Controller [19e5:a239]
|
|      \-02.0  Huawei Technologies Co., Ltd. HiSilicon USB 3.0 Host Controller [19e5:a238]
+-[0000:74]---01.0-[76]---
|
|      +-02.0  Huawei Technologies Co., Ltd. HiSilicon SAS 3.0 HBA [19e5:a230]
|
|      +-03.0  Huawei Technologies Co., Ltd. HiSilicon AHCI HBA [19e5:a235]
|
|      \-04.0  Huawei Technologies Co., Ltd. HiSilicon SAS 3.0 HBA [19e5:a230]
\-[0000:00]---00.0-[01]---
|
|      +-04.0-[02]---
|
|      +-08.0-[03]----00.0  LSI Logic / Symbios Logic MegaRAID Tri-Mode SAS3408 [1000:0017]
|
|      +-0c.0-[04-09]----00.0-[05-09]---+-00.0-[06]----00.0  Huawei Technologies Co., Ltd. Hi1822 Family (4*25GE) [19e5:1822]
|
|      |      +-01.0-[07]----00.0  Huawei Technologies Co., Ltd. Hi1822 Family (4*25GE) [19e5:1822]
|
|      |      +-02.0-[08]----00.0  Huawei Technologies Co., Ltd. Hi1822 Family (4*25GE) [19e5:1822]
|
|      |      \-03.0-[09]----00.0  Huawei Technologies Co., Ltd. Hi1822 Family (4*25GE) [19e5:1822]
|
|      +-10.0-[0a]----00.0  Huawei Technologies Co., Ltd. iBMA Virtual Network Adapter [19e5:1710]
|
|      +-11.0-[0b]----00.0  Huawei Technologies Co., Ltd. Hi1710 [iBMC Intelligent Management system chip w/VGA support] [19e5:1711]
|
|      \-12.0-[0c]---
```

PCI域: 总线号(8bit): 设备号(5bit): 功能号(3bit)

例: 0000: 01: 00.0

# 查看系统PCI设备

- 查看pci设备拓扑: `lspci -tv (-Dnn)`
- 查看配置空间: `lspci -vvv -s BDF号`
- 以二进制方式查看配置空间: `lspci -xxx -s BDF号 (-xxxx 查看扩展配置空间)`

```
[root@server1 ~]# lspci -vvv -s 08:00.0
08:00.0 Ethernet controller: Huawei Technologies Co., Ltd. Hi1822 Family (4*25GE) (rev 45)
Subsystem: Huawei Technologies Co., Ltd. Device d136
Control: I/O- Mem+ BusMaster+ SpecCycle- MemWINV- VGASnoop- ParErr+ Stepping- SERR+ FastB2B- DisINTx+
Status: Cap+ 66MHz- UDF- FastB2B- ParErr- DEVSEL=fast >TAbort- <TAbort- <MAbort- >SERR- <PERR- INTx+
Latency: 0, Cache Line Size: 32 bytes
NUMA node: 0
Region 0: Memory at 80003400000 (64-bit, prefetchable) [size=128K]
Region 2: Memory at 800034e0000 (64-bit, prefetchable) [size=32K]
Region 4: Memory at 80003300000 (64-bit, prefetchable) [size=1M]
Capabilities: [40] Express (v2) Endpoint, MSI 00
    DevCap: MaxPayload 512 bytes, PhantFunc 0, Latency L0s unlimited, L1 unlimited
              ExtTag+ AttnBtn- AttnInd- PwrInd- RBE+ FLReset+ SlotPowerLimit 0.000W
    DevCtl: CorrErr+ NonFatalErr+ FatalErr+ UnsupReq+
              RlxdOrd+ ExtTag+ PhantFunc- AuxPwr- NoSnoop+ FLReset-
              MaxPayload 256 bytes, MaxReadReq 512 bytes
    DevSta: CorrErr+ NonFatalErr- FatalErr- UnsupReq+ AuxPwr+ TransPend-
    LnkCap: Port #0, Speed 8GT/s, Width x16, ASPM not supported
              ClockPM- Surprise- LLActRep- BwNot- ASPMOptComp+
    LnkCtl: ASPM Disabled; RCB 128 bytes Disabled- CommClk-
              ExtSynch- ClockPM- AutWidDis- BWInt- AutBWInt-
    LnkSta: Speed 8GT/s (ok), Width x16 (ok)
              TrErr- Train- SlotClk- DLActive- BWMgmt- ABWMgmt-
    DevCap2: Completion Timeout: Range B, TimeoutDis+, LTR-, OBFF Not Supported
              AtomicOpsCap: 32bit+ 64bit+ 128bitCAS+
    DevCtl2: Completion Timeout: 50us to 50ms, TimeoutDis-, LTR-, OBFF Disabled
              AtomicOpsCtl: ReqEn-
    LnkCtl2: Target Link Speed: 8GT/s, EnterCompliance- SpeedDis-
              Transmit Margin: Normal Operating Range, EnterModifiedCompliance- ComplianceSOS-
              Compliance De-emphasis: -6dB
    LnkSta2: Current De-emphasis Level: -3.5dB, EqualizationComplete+, EqualizationPhase1+
              EqualizationPhase2+, EqualizationPhase3+, LinkEqualizationRequest-
Capabilities: [80] MSI: Enable- Count=1/32 Maskable+ 64bit+
    Address: 0000000000000000 Data: 0000
    Masking: 00000000 Pending: 00000000
Capabilities: [a0] MSI-X: Enable+ Count=128 Masked-
    Vector table: BAR=2 offset=00000000
```



# 查看系统PCI设备

- 查看pci设备拓扑: `lspci -tv (-Dnn)`
- 查看配置空间: `lspci -vvv -s BDF号`
- 以二进制方式查看配置空间: `lspci -xxx -s BDF号 (-xxxx 查看扩展配置空间)`

```
[root@server1 ~]# lspci -xxx -s 08:00.0
08:00.0 Ethernet controller: Huawei Technologies Co., Ltd. Hi1822
00: e5 19 22 18 46 05 18 00 45 00 00 02 08 00 00 00
10: 0c 00 40 03 00 08 00 00 0c 00 4e 03 00 08 00 00
20: 0c 00 30 03 00 08 00 00 00 00 00 00 e5 19 36 d1
30: 00 00 00 00 40 00 00 00 00 00 00 00 ff 00 00 00
40: 10 80 02 00 e2 8f 00 10 3f 29 19 00 03 f1 43 00
50: 08 00 03 01 00 00 00 00 00 00 00 00 00 00 00 00
60: 00 00 00 00 92 03 00 00 00 00 00 00 00 0e 00 00
70: 03 00 1f 00 00 00 00 00 00 00 00 00 00 00 00 00
80: 05 a0 8a 01 00 00 00 00 00 00 00 00 00 00 00 00
90: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
a0: 11 b0 7f 80 02 00 00 00 02 40 00 00 00 00 00 00
b0: 01 c0 03 f8 00 00 00 00 00 00 00 00 00 00 00 00
c0: 03 00 28 80 30 78 ff ff 00 00 00 00 00 00 00 00
d0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
e0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
f0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
```

怎样唯一标识一个PCI设备？

怎样访问一个PCI设备，设备驱动  
ioremap哪个地址呢？

# PCI配置空间

## ■ 配置空间

- 可以直接访问
- 用于发现设备
- 用于分配mem/io地址

## ■ MEM空间

- 外设驱动访问

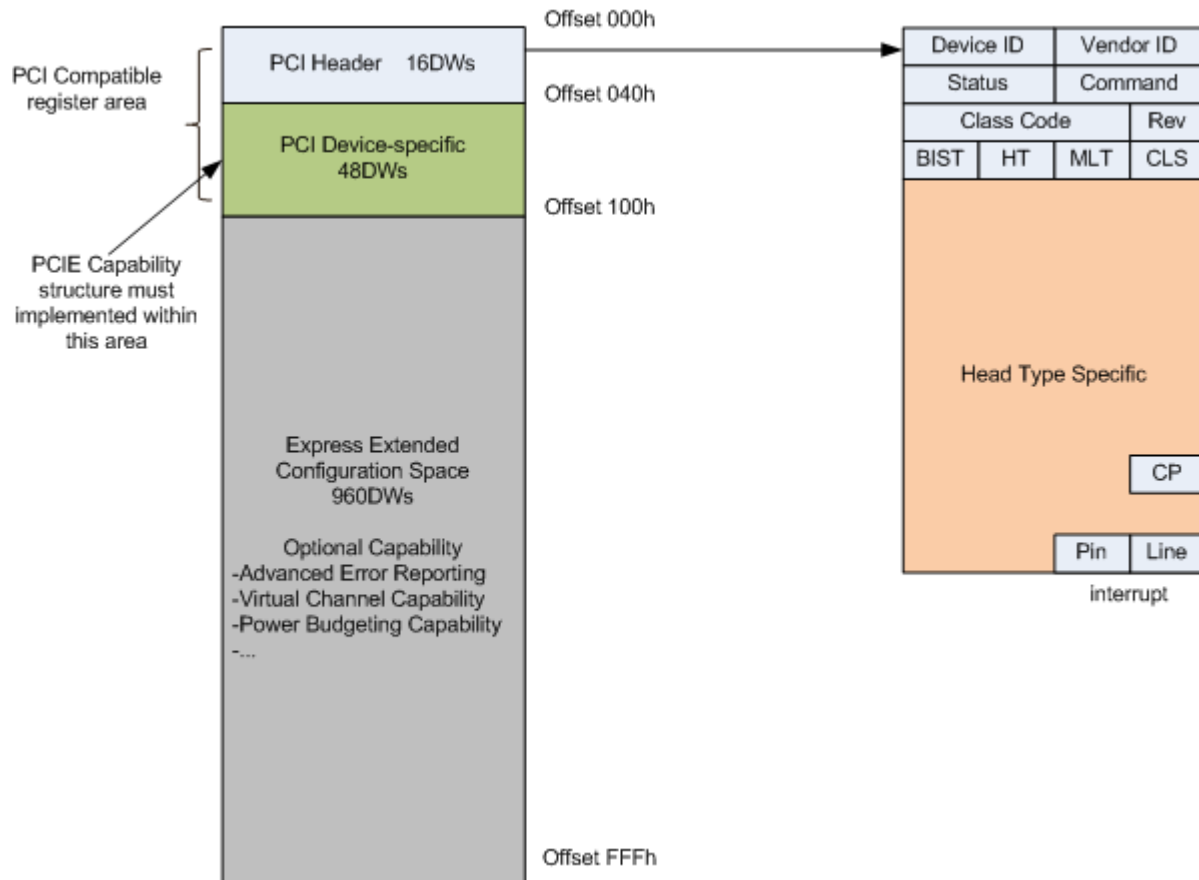
## ■ IO空间

- 外设驱动访问



访问配置空间

- IO(CFC/CF8寄存器)
- ECAM方式



基本配置空间256字节，扩展配置空间4K

# PCI配置空间

64  
Bytes

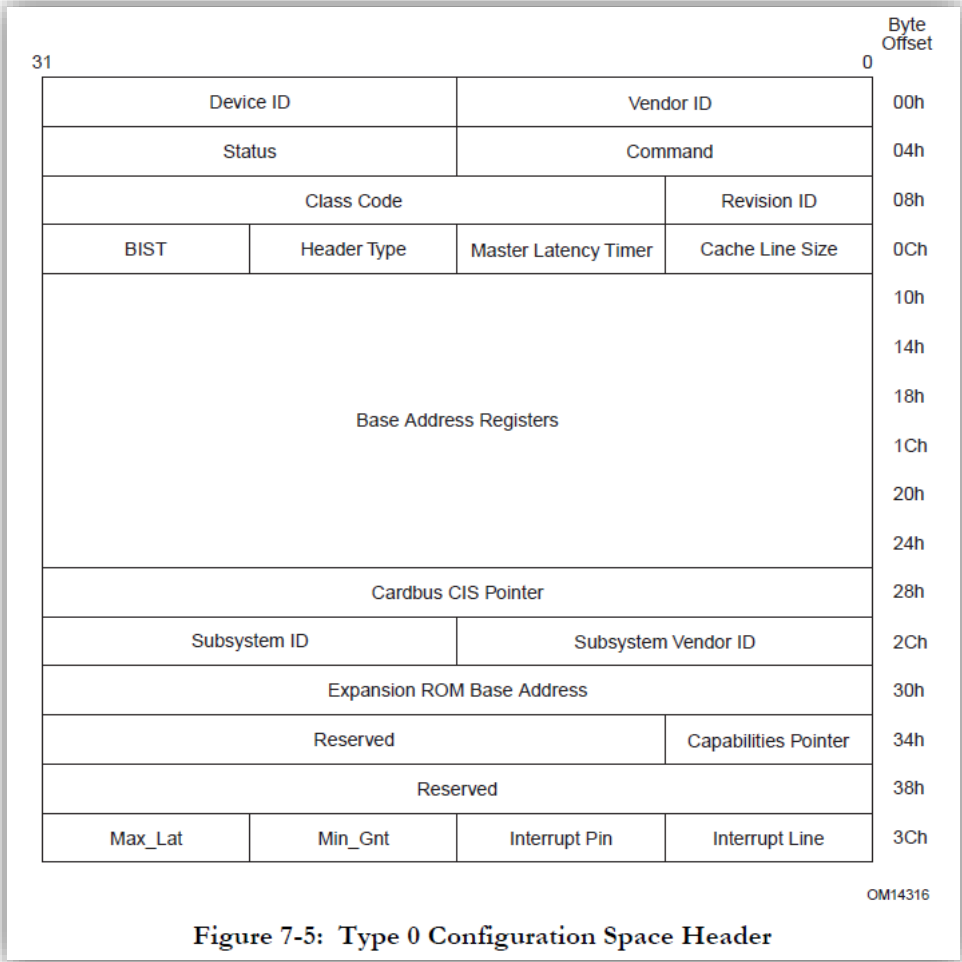


Figure 7-5: Type 0 Configuration Space Header

EndPoint 配置空间头

**Device ID:** 生产厂商指定的16位硬件设备编码

**Vendor ID:** 标识硬件制造商，如Intel 0x8086

**Status:** 设备相关状态位寄存器；

**Commad:** 设备控制寄存器，包括各种使能

**Class:** 设备所属类编码

**Header Type:** 设备所属PCI类型，端口设备或桥设备

**Base :** 基地址寄存器(min memory 128bytes)

**Interrupt Pin:** 对应PCI INTA#-INTD#四个中断引脚

**Interrupt Line:** 保存中断路由信息，对应中断控制器的引脚(IRQ0-IRQ15)

**Subsystem vendorId, subsystem deviceId:** 进一步识别设备字段

**Expansion ROM Base Register:** 设备本身提供的Firmware映射信息的地址寄存器



# PCI配置空间

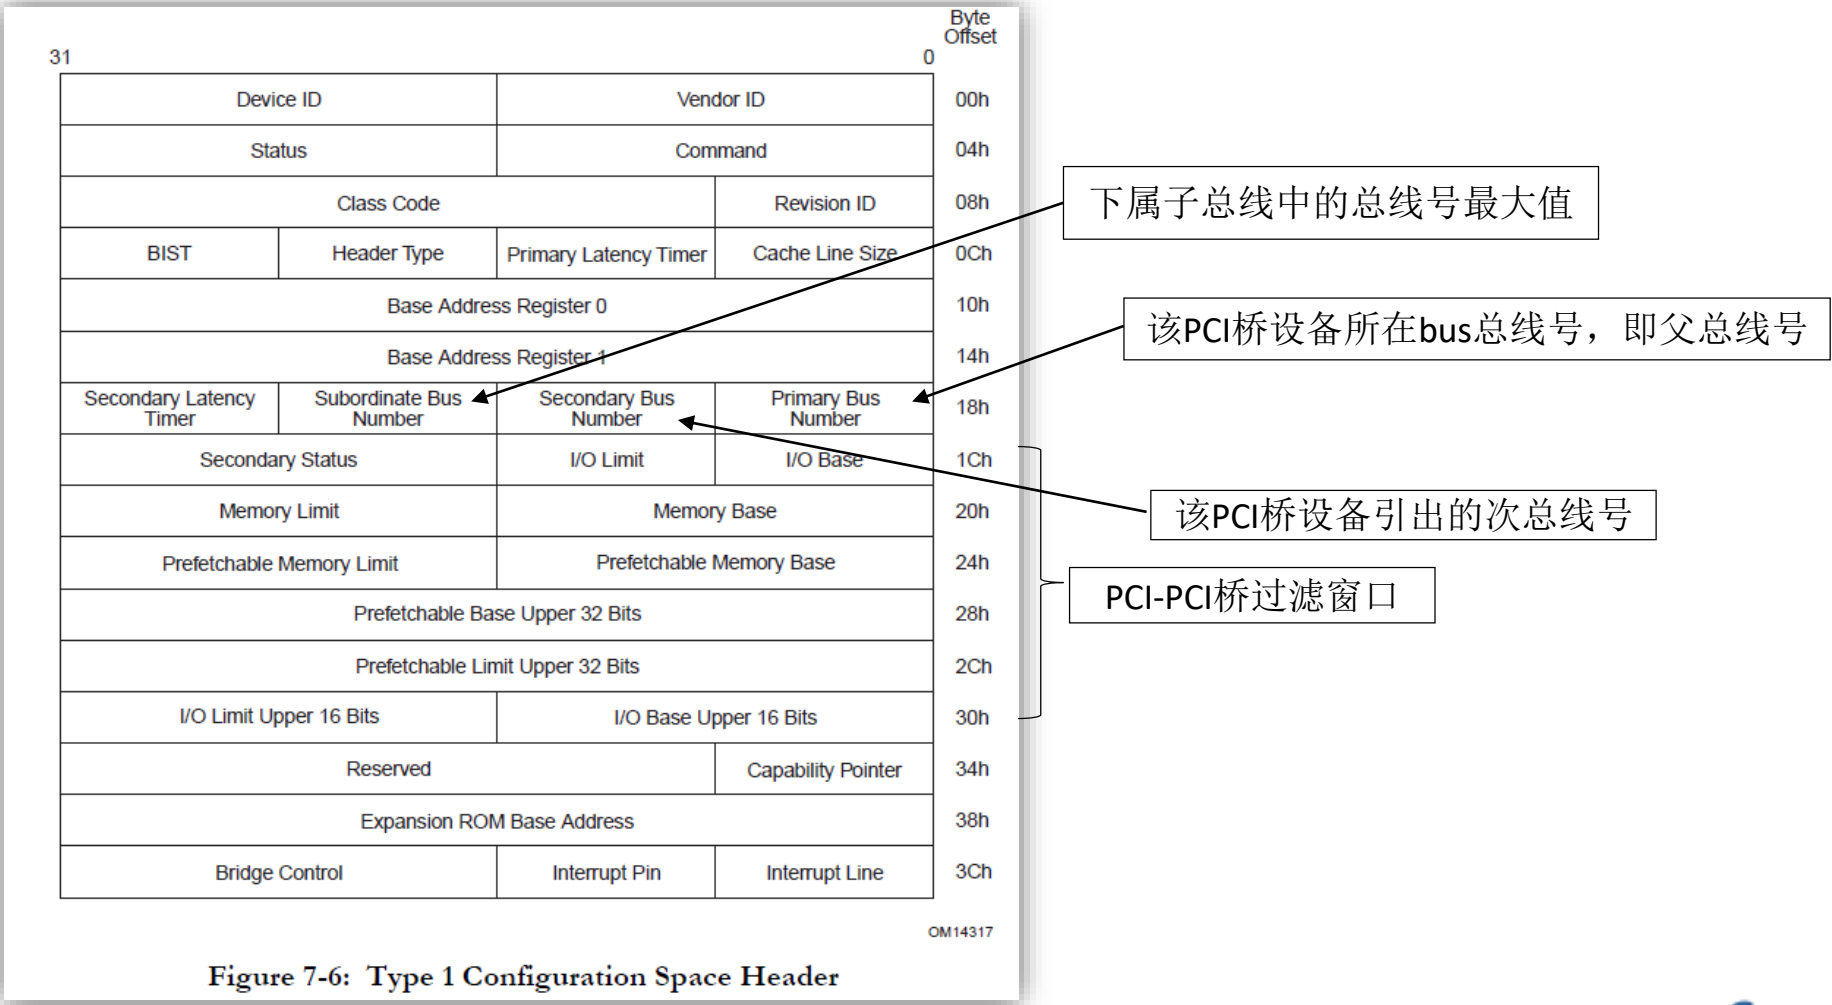



Figure 7-6: Type 1 Configuration Space Header

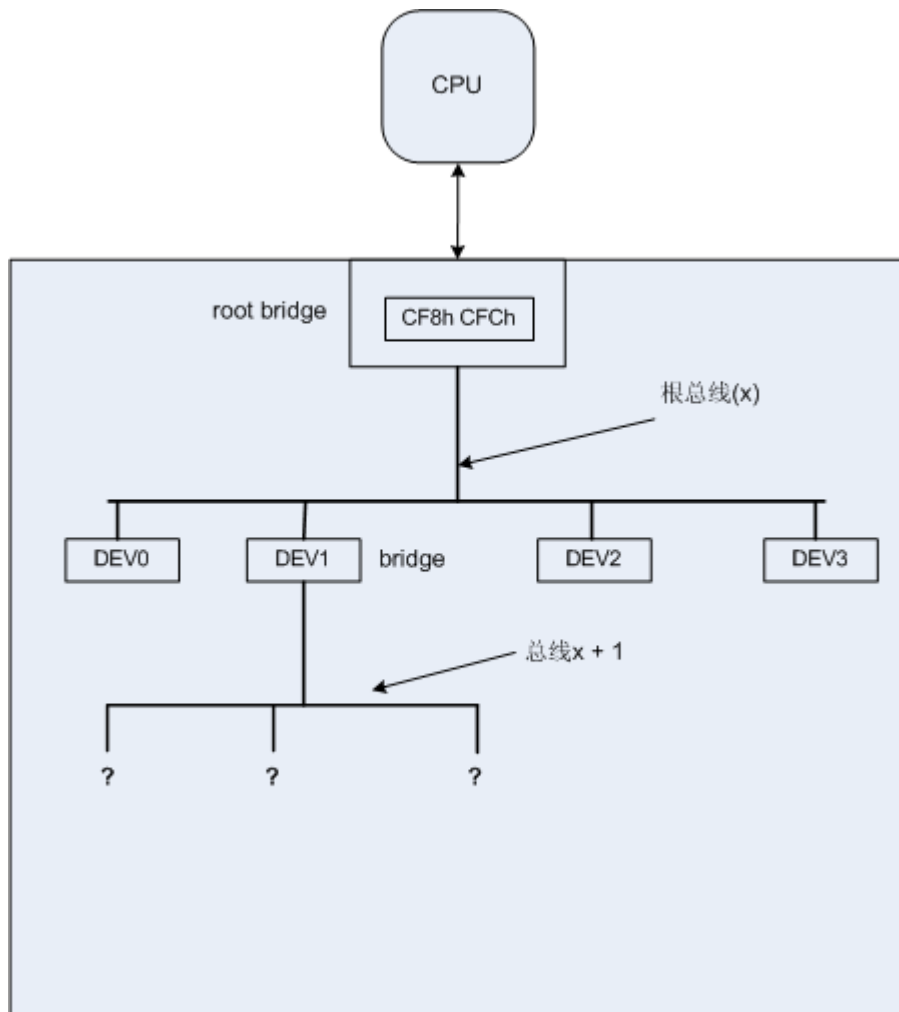
桥片配置空间头

# PCI初始化——设备枚举



- 遍历0号总线下的所有设备号功能号，device id 和vendor id非全0或全f即存在设备
- 建立PCI设备对应的内核结构struct pci\_dev

# PCI初始化——设备枚举



- 遍历0号总线下的所有设备号功能号，device id 和vendor id非全0或全f即存在设备
- 建立PCI设备对应的内核结构struct pci\_dev
- 如果某个设备为桥设备则引出1号总线
- 遍历1号总线下的所有设备号功能号，找出存在的PCI设备
- 依次递归



# PCI初始化代码分析

```
static const struct acpi_device_id root_device_ids[] = {
    {"PNP0A03", 0},
    {"", 0},
};

static struct acpi_scan_handler pci_root_handler = {
    .ids = root_device_ids,
    .attach = acpi_pci_root_add,
    .detach = acpi_pci_root_remove,
    .hotplug = {
        .enabled = true,
        .scan_dependent = acpi_pci_root_scan_dependent,
    },
};
```

acpi\_pci\_root\_add() //drivers/acpi/pci\_root.c

->pci\_acpi\_scan\_root() //arch/arm64/kernel/pci.c

->acpi\_pci\_root\_create();

->pci\_scan\_child\_bus(bus); //枚举设备

->pci\_assign\_unassigned\_root\_bus\_resources(bus); //分配资源

```
Device (PCI2)
{
    Name (_HID, "PNP0A08") // PCI Express Root Bridge
    Name (_CID, "PNP0A03") // Compatible PCI Root Bridge
    Name(_SEG, 2) // Segment of this Root complex
    Name(_BBN, 0x80) // Base Bus Number
    Name(_CCA, 1)
    Method (_CRS, 0, Serialized) { // Root complex resources
        Name (RBUF, ResourceTemplate () {
            WordBusNumber ( // Bus numbers assigned to this root
                ResourceProducer, MinFixed, MaxFixed, PosDecode,
                0, // AddressGranularity
                0x80, // AddressMinimum - Minimum Bus Number
                0x87, // AddressMaximum - Maximum Bus Number
                0, // AddressTranslation - Set to 0
                0x8 // RangeLength - Number of Busses
            )
            QWordMemory ( // 64-bit BAR Windows
                ResourceProducer,
                PosDecode,
                MinFixed,
                MaxFixed,
                Cacheable,
                ReadWrite,
                0x0, // Granularity
                0xa8800000, // Min Base Address
                0xfffffff, // Max Base Address
                0x0, // Translate
                0x77f0000 // Length
            )
            QWordIO (
                ResourceProducer,
                MinFixed,
                MaxFixed,
                PosDecode,
                EntireRange,
                0x0, // Granularity
                0x0, // Min Base Address
                0xffff, // Max Base Address
                0xffff0000, // Translate
                0x10000 // Length
            )
        }) // Name (RBUF)
        Return (RBUF)
    } // Method (_CRS)
```

# PCI设备枚举代码分析

```
static unsigned int pci_scan_child_bus_extend(struct pci_bus *bus,
/* ◀ ▶ ⏪ ⏩ ⏴ ⏵ ⏶ ⏷ ? */
                                unsigned int available_buses)
{
    unsigned int used_buses, normal_bridges = 0, hotplug_bridges = 0;
    unsigned int start = bus->busn_res.start;
    unsigned int devfn, fn, cmax, max = start;
    struct pci_dev *dev;
    int nr_devs;

    dev_dbg(&bus->dev, "scanning bus\n");

    /* Go find them, Rover! */
    for (devfn = 0; devfn < 256; devfn += 8) {
        nr_devs = pci_scan_slot(bus, devfn);

        /*
         * The Jailhouse hypervisor may pass individual functions of a
         * multi-function device to a guest without passing function 0.
         * Look for them as well.
         */
        if (jailhouse_naravirt() && nr_devs == 0) {

/*
 * Scan bridges that are already configured. We don't touch them
 * unless they are misconfigured (which will be done in the second
 * scan below).
 */
        for each_pci_bridge(dev, bus) {
            cmax = max;
            max = pci_scan_bridge_extend(bus, dev, max, 0, 0);

            /*
             * Reserve one bus for each bridge now to avoid extending
             * hotplug bridges too much during the second scan below.
             */
            used_buses++;
            if (cmax - max > 1)
                used_buses += cmax - max - 1;
        }
    }
}
```

```
int pci_scan_slot(struct pci_bus *bus, int devfn)
/* ◀ ▶ ⏪ ⏩ ⏴ ⏵ 🏠 ? */
{
    unsigned fn, nr = 0;
    struct pci_dev *dev;

    if (only_one_child(bus) && (devfn > 0))
        return 0; /* Already scanned the entire slot */

    dev = pci_scan_single_device(bus, devfn);
    if (!dev)
        return 0;
    if (!pci_dev_is_added(dev))
        nr++;

    for (fn = next_fn(bus, dev, 0); fn > 0; fn = next_fn(bus, dev, fn)) {
        dev = pci_scan_single_device(bus, devfn + fn);
        if (dev) {
            if (!pci_dev_is_added(dev))
                nr++;
            dev->multifunction = 1;
        }
    }
}
```



# PCI设备对应的内核结构体

```
struct pci_bus {
    struct list_head node;           /* Node in list of buses */
    struct pci_bus *parent;          /* Parent bus this bridge is on */
    struct list_head children;       /* List of child buses */
    struct list_head devices;        /* List of devices on this bus */
    struct pci_dev *self;            /* Bridge device as seen by parent */
    struct list_head slots;          /* List of slots on this bus;
                                     protected by pci_slot_mutex */

    struct resource *resource[PCI_BRIDGE_RESOURCE_NUM];
    struct list_head resources;      /* Address space routed to this bus */
    struct resource busn_res;        /* Bus numbers routed to this bus */

    struct pci_ops *ops;             /* Configuration access functions */
    struct pci_ops *backup_ops;
    struct msi_controller *msi;      /* MSI controller */
    void *sysdata;                  /* Hook for sys-specific extension */
    struct proc_dir_entry *procdir; /* Directory entry in /proc/bus/pci */

    unsigned char number;            /* Bus number */
    unsigned char primary;           /* Number of primary bridge */
};
```

PCI 总线对应的内核结构体

```
/* The pci_dev structure describes PCI devices */
struct pci_dev {
    struct list_head bus_list;      /* Node in per-bus list */
    struct pci_bus *bus;            /* Bus this device is on */
    struct pci_bus *subordinate;    /* Bus this device bridges to */

    void *sysdata;                  /* Hook for sys-specific extension */
    struct proc_dir_entry *procent; /* Device entry in /proc/bus/pci */
    struct pci_slot *slot;          /* Physical slot this device is in */

    unsigned int devfn;             /* Encoded device & function index */
    unsigned short vendor;
    unsigned short device;
    unsigned short subsystem_vendor;
    unsigned short subsystem_device;
    unsigned int class;             /* 3 bytes: (base, sub, prog-if) */
    u8 revision;                   /* PCI revision, low byte of class word */
    u8 hdr_type;                   /* PCI header type ('multi' flag masked out) */

#ifdef CONFIG_PCIEAER
    u16 aer_cap;                   /* AER capability offset */
    struct aer_stats *aer_stats;    /* AER stats for this device */
#endif

    u8 pcie_cap;                   /* PCIe capability offset */
    u8 msi_cap;                    /* MSI capability offset */
    u8 msix_cap;                   /* MSI-X capability offset */
    u8 pcie_mps:3;                 /* PCIe Max Payload Size Supported */
    u8 rom_base_reg;               /* Config register controlling ROM */
    u8 pin;                        /* Interrupt pin this device uses */
    u16 pcie_flags_reg;            /* Cached PCIe Capabilities Register */
    unsigned long *dma_alias_mask; /* Mask of enabled devfn aliases */

    struct pci_driver *driver;      /* Driver bound to this device */
};

struct device dev;                /* Generic device interface */

int cfg_size;                     /* Size of config space */

/*
 * Instead of touching interrupt line and base address registers
 * directly, use the values stored here. They might be different!
 */
unsigned int irq;
struct resource resource[DEVICE_COUNT_RESOURCE]; /* I/O and memory regions + expansion ROMs */
```

PCI设备对应的内核结构体



# PCI设备驱动

```
int pci_setup_device(struct pci_dev *dev)
/* ◀ ▶ ◀ ▶ ⏏ ⏏ ⏏ ? */
{
    u32 class;
    u16 cmd;
    u8 hdr type;
    int pos = 0;
    struct pci_bus region;
    struct resource *res;

    hdr type = pci_hdr_type(dev);

    dev->sysdata = dev->bus->sysdata;
    dev->dev.parent = dev->bus->bridge;
    dev->dev.bus = &pci_bus_type;
    dev->hdr_type = hdr type & 0x7f;
    dev->multifunction = !(hdr type & 0x80);
    dev->error_state = pci_channel_io_normal;
    set_pcie_port_type(dev);

    struct bus_type pci_bus_type = {
        .name           = "pci",
        .match           = pci_bus_match,
        .uevent          = pci_uevent,
        .probe           = pci_device_probe,
        .remove          = pci_device_remove,
        .shutdown        = pci_device_shutdown,
        .dev_groups      = pci_dev_groups,
        .bus_groups      = pci_bus_groups,
        .drv_groups      = pci_drv_groups,
        .pm              = PCI_PM_OPS_PTR,
        .num_vf          = pci_bus_num_vf,
        .dma_configure    = pci_dma_configure,
    };
    EXPORT_SYMBOL(pci_bus_type);
}
```

PCI驱动

```
static struct pci_driver sec_pci_driver = {
    .name = "hisi_sec2",
    .id_table = sec_dev_ids,
    .probe = sec_probe,
    .remove = sec_remove,
    .err_handler = &sec_err_handler,
    .sriov_configure = sec_sriov_configure,
    .shutdown = hisi_qm_dev_shutdown,
};
```

```
static int sec_probe(struct pci_dev *pdev, const struct pci_device_id *id)
/* ◀ ▶ ◀ ▶ ⏏ ⏏ ⏏ ? */
{
    struct sec_dev *sec;
    struct hisi_qm *qm;
    int ret;

    sec = devm_kzalloc(&pdev->dev, sizeof(*sec), GFP_KERNEL);
    if (!sec)
        return -ENOMEM;

#ifdef CONFIG_CRYPTO_QM_UACCE
    qm->phys_base = pci_resource_start(pdev, PCI_BAR_2);
    qm->size = pci_resource_len(qm->pdev, PCI_BAR_2);
#endif

    qm->io_base = devm_ioremap(pdev, pci_resource_start(pdev, PCI_BAR_2),
                             pci_resource_len(qm->pdev, PCI_BAR_2));
    if (!qm->io_base) {
        ret = -EIO;
        goto err_release_mem_regions;
    }

    ret = dma_set_mask_and_coherent(dev, DMA_BIT_MASK(64));
    if (ret < 0) {
        dev_err(dev, "Failed to set 64 bit dma mask %d", ret);
        goto err_iounmap;
    }
    pci_set_master(pdev);

    num_vec = qm->ops->get_irq_num(qm);
    ret = pci_alloc_irq_vectors(pdev, num_vec, num_vec, PCI_IRQ_MSI);
    if (ret < 0) {
        dev_err(dev, "Failed to enable MSI vectors!\n");
        goto err_iounmap;
    }
}
```

PCI总线

## ✓ openEuler kernel gitee 仓库

源代码仓库

<https://gitee.com/openeuler/kernel>

欢迎大家多多 Star，多多参与社区开发，多多贡献补丁。

## ✓ maillist、issue、bugzilla

可以通过邮件列表、issue、bugzilla 参与社区讨论

欢迎大家多多讨论问题，发现问题多提 issue、bugzilla

<https://gitee.com/openeuler/kernel/issues>

<https://bugzilla.openeuler.org>

[kernel@openeuler.org](mailto:kernel@openeuler.org)


## ✓ openEuler kernel SIG 微信技术交流群

请扫描右方二维码添加小助手微信

或者直接添加小助手微信（微信号：openeuler-kernel）

备注“交流群”或“技术交流”

加入 openEuler kernel SIG 技术交流群



技术交流



# Thank you